

## 18. Monte Carlo Methods

### Monte Carlo

#### The Monte Carlo method

The Monte Carlo method is such an important pricing technique that we dedicate the entire chapter to this. Originating from a sampling technique developed in statistical physics, Monte Carlo methods are now widely used for pricing derivatives, especially when dealing with complex instruments or market models where analytical solutions are difficult or impossible to obtain.

The basic idea behind Monte Carlo simulation is to use randomness to solve problems that might be deterministic in principle. In the context of pricing derivatives, Monte Carlo methods involve simulating a large number of random paths for the underlying asset price to estimate the expected value of a derivative's payoff.

The biggest advantage is that the Monte Carlo engine can be written once and updating or adding new SDE's is relatively easy. The same can be said for implementing the payoff of a derivative. Compared to finite difference methods or other methods that focus on solving the PDE, the implementation can be done agnostic of the product.

#### Key Concepts of Monte Carlo Simulation

1. **Random Sampling:** Monte Carlo relies on generating random samples from a probability distribution. For derivative pricing, these samples represent possible future paths of the underlying asset price based on a specified stochastic process (like Geometric Brownian Motion).
2. **Expectation Estimation:** The price of a derivative is often the expected value of its future payoff under the risk-neutral measure. Monte Carlo simulates many scenarios (paths) of the underlying asset's future prices and calculates the payoff for each scenario. By averaging these payoffs, it estimates the expected value.
3. **Law of Large Numbers:** This fundamental statistical principle states that as the number of trials or simulations increases, the average of the results will converge to the expected value. Monte Carlo leverages this by using a large number of simulations to get a more accurate estimate of the derivative's price.
4. **Versatility and Flexibility:** Monte Carlo methods can be applied to a wide range of problems, especially those that are difficult to solve analytically. They are particularly useful for pricing derivatives with complex features, such as path dependence (e.g., Asian options) or early exercise (e.g., American options).

#### Illustration of Monte Carlo in Derivative Pricing

Let's say we want to price a European call option, which gives the right to buy a stock at a strike price  $K$  at a future date  $T$ .

- **Step 1:** Define the model for the stock price, typically using a stochastic process like Geometric Brownian Motion.
- **Step 2:** Simulate a large number of possible future stock price paths up to the option's expiration date.
- **Step 3:** Calculate the payoff for each path (e.g.,  $\max(S_T - K, 0)$  for a call option).
- **Step 4:** Average these payoffs across all simulated paths.
- **Step 5:** Discount this average payoff to present value using the risk-free rate to get the estimated option price.

## Steps for Monte Carlo Simulation

Here's a step-by-step breakdown of how Monte Carlo simulation is typically applied to price a derivative:

### Step 1: Model the Stochastic Process of the Underlying Asset

The first step is to define a model for the stochastic process that governs the underlying asset price. One common choice is the **Geometric Brownian Motion (GBM)**, defined by the stochastic differential equation (SDE):

$$dS_t = S_t(\mu dt + \sigma dW_t)$$

where:

- $S_t$  is the asset price at time  $t$ ,
- $\mu$  is the drift rate of the asset (often set to the risk-free rate  $r$  in a risk-neutral framework),
- $\sigma$  is the volatility of the asset,
- $W_t$  is a Wiener process or Brownian motion.

In a risk-neutral framework, we typically assume that the drift  $\mu = r - q$ , where  $r$  is the risk-free rate, and  $q$  is the continuous dividend yield.

### Step 2: Discretize the Stochastic Process

To simulate the paths of the underlying asset, we discretize the continuous-time model over small time steps  $\Delta t$ . Divide the whole time interval (from  $t_0$  to the maturity  $T$  of the option) into a number ( $n$ ) periods such that  $\Delta t = T/n$ .

A simple discretization of the GBM model using the Euler scheme (see [Discretisation of SDEs](#)) is:

$$S_{t+\Delta t} = S_t \left( 1 + (r - q)\Delta t + \sigma\sqrt{\Delta t} Z \right)$$

where:

- $Z$  is a standard normal random variable,  $Z \sim N(0, 1)$ .

Practically speaking, this means, start at  $t_0$  with  $S(t_0) = S_0$ . Then for each time step, pick a random number  $Z$  and calculate the asset value at the end of the interval, using the above discretization formula until the maturity is reached and  $S(T)$  is simulated.

### Step 3: Simulate Asset Price Paths

Next, generate a large number of simulated paths (say,  $N$  paths) of the underlying asset price from time 0 to the option's maturity  $T$ . For each path, start with the initial asset price  $S_0$  and iteratively apply the discretization formula to get  $S_t$  at each time step.

Typically, the number of simulations has to be large. The exact requirement depends on the complexity of the option, but a rough guideline is at least 100,000 and as much as the computing power allows.

### Step 4: Calculate Payoffs for Each Path

For each simulated path, calculate the payoff of the derivative. The payoff depends on the type of derivative. For example:

- **European Call Option:**  $\max(S_T - K, 0)$
- **European Put Option:**  $\max(K - S_T, 0)$

where  $K$  is the strike price and  $S_T$  is the simulated price of the underlying asset at maturity  $T$ .

### Step 5: Average the Payoffs and Discount to Present Value

Compute the average payoff across all the simulated paths. To obtain the present value of the derivative, discount the average payoff back to today using the risk-free rate  $r$ :

$$\text{Price} = e^{-rT} \frac{1}{N} \sum_{i=1}^N \text{Payoff}_i$$

where  $\text{Payoff}_i$  is the payoff of the  $i$ -th simulated path.

## Advantages and Limitations of Monte Carlo

### Advantages:

- **Flexibility:** Monte Carlo can handle a wide range of derivatives and models.
- **Path Dependence:** It is well-suited for derivatives whose payoff depends on the entire path of the underlying asset price (eg. Asian options, barrier options, lookback options, etc)
- The **greeks** can also be estimated through a Monte Carlo.

### Limitations:

- **Computationally Intensive:** Requires a large number of simulations for accurate pricing, which can be time-consuming.
- **Convergence:** Convergence to the correct price can be slow, especially for options with complex payoffs.
- American-style options do not fall under this framework. There are numerical work-arounds for these but these are more complex to implement.

## Discretization process for the SDE

### Discretisation of SDEs

Discretizing a Stochastic Differential Equation (SDE) is a crucial step in using numerical methods to simulate stochastic processes, such as those involved in financial modeling and

derivative pricing. Discretization involves approximating the continuous-time SDE by a discrete-time process, which can be simulated on a computer. Here's a comprehensive overview of how to discretize an SDE.

## Understanding the SDE

A typical SDE that models the dynamics of an asset price  $S_t$  can be written as:

$$dS_t = \mu(S_t, t) dt + \sigma(S_t, t) dW_t$$

where:

- $S_t$  is the asset price at time  $t$ .
- $\mu(S_t, t)$  is the drift term, representing the deterministic part of the asset's return.
- $\sigma(S_t, t)$  is the volatility term, representing the random part of the asset's return.
- $dW_t$  is an increment of a Wiener process (or Brownian motion), which is normally distributed with mean 0 and variance  $dt$ ,  $dW_t \sim \mathcal{N}(0, dt)$ .

## Purpose of Discretization

Discretization converts the continuous-time SDE into a form that can be implemented in a discrete-time simulation. This involves approximating the continuous changes over an infinitesimal interval  $dt$  with discrete changes over a finite time step  $\Delta t$ .

## Discretization Methods

Several methods exist for discretizing SDEs, each with varying levels of accuracy and complexity. Here are some commonly used methods:

### Euler Method

The Euler method is the simplest and most widely used method for discretizing an SDE. It is a straightforward extension of the Euler method for ordinary differential equations (ODEs) to include stochastic terms.

Given an SDE:

$$dS_t = \mu(S_t, t) dt + \sigma(S_t, t) dW_t$$

The Euler discretization is:

$$S_{t+\Delta t} = S_t + \mu(S_t, t)\Delta t + \sigma(S_t, t)\sqrt{\Delta t} Z$$

where:

- $\Delta t$  is the discrete time step.
- $Z$  is a standard normal random variable,  $Z \sim \mathcal{N}(0, 1)$ .

### Advantages:

- Simple to implement.
- Computationally efficient.

### Limitations:

- The Euler method is only a first-order method, meaning its error is proportional to  $\Delta t$ . It may not be accurate enough for processes requiring fine precision.

## Milstein Method

The Milstein method improves upon the Euler method by including an additional term that accounts for the variation in volatility. For SDEs with non-constant volatility, the Milstein method provides a more accurate approximation.

The Milstein discretization is:

$$S_{t+\Delta t} = S_t + \mu(S_t, t)\Delta t + \sigma(S_t, t)\sqrt{\Delta t}Z + \frac{1}{2}\sigma(S_t, t)\sigma'(S_t, t)(\Delta W_t^2 - \Delta t)$$

where  $\sigma'(S_t, t)$  denotes the partial derivative of  $\sigma(S_t, t)$  with respect to  $S_t$ , and  $\Delta W_t = \sqrt{\Delta t}Z$ .

### Advantages:

- More accurate than Euler for SDEs with variable volatility.
- Second-order accurate in terms of  $\Delta t$ .

### Limitations:

- Slightly more complex to implement due to the additional term.

## Higher-Order Methods (e.g., Runge-Kutta Methods for SDEs)

Higher-order methods, such as stochastic Runge-Kutta methods, offer greater accuracy by using more terms from the Taylor expansion of the solution to the SDE. These methods are particularly useful for SDEs where high accuracy is required over long time horizons.

The general form involves multiple evaluations of drift and diffusion functions over each time step, weighted by coefficients derived from the desired order of accuracy.

### Advantages:

- High accuracy, especially for small  $\Delta t$ .

### Limitations:

- Computationally intensive.
- More complex to implement.

## Transformations

Sometimes it is possible to transform the SDE to a simpler form. For example, for the Geometric Brownian Motion, the SDE is:

$$dS_t = \mu S_t dt + \sigma S_t dW_t$$

we can transform as  $Y_t = \ln(S_t)$  and from Ito's lemma we find that

$$dY_t = (\mu - \frac{1}{2}\sigma^2)dt + \sigma dW_t$$

which we can then discretize with the Euler method to find:

$$Y_{t+\Delta t} = (\mu - \frac{1}{2}\sigma^2)\Delta t + \sigma\sqrt{\Delta t}Z$$

and if we rewrite this in terms of the original price process  $S_t$ , we obtain

$$S_{t+\Delta t} = S_t \exp \left( \left( \mu - \frac{1}{2} \sigma^2 \right) \Delta t + \sigma \sqrt{\Delta t} Z \right)$$

We can immediately see that one of the main advantages is that  $S_t > 0$  by construction. In the Euler scheme, this is not necessarily guaranteed.

## Choosing a Discretization Method

The choice of discretization method depends on:

- **Accuracy Requirements:** Higher-order methods provide better accuracy but at the cost of increased computational complexity.
- **Computational Resources:** Simpler methods like Euler are faster but less accurate.
- **Properties of the SDE:** If the volatility term  $\sigma(S_t, t)$  is not constant, methods like the Milstein method may be preferable.

When implementing SDE discretization:

- **Time Step  $\Delta t$ :** Choose a sufficiently small time step to balance accuracy and computational cost. Smaller time steps reduce discretization error but increase the number of computations.
- **Random Number Generation:** Ensure high-quality random number generation for the standard normal variable  $Z$ .
- **Error Analysis:** Evaluate the convergence and error of the discretization method by comparing results for decreasing time steps.

## Random number generation

Generating standard random numbers, particularly **standard normal random variables** (i.e., variables that follow a normal distribution with a mean of 0 and a standard deviation of 1), is essential in many numerical methods, including Monte Carlo simulations and stochastic processes like those in financial modeling. Here are the most common methods for generating these random numbers:

### Random Number Generation: Overview

**Random number generation** is crucial in various applications, from simulations and statistical sampling to cryptography and gaming. Generally, random numbers can be generated by hardware-based systems (true random numbers) or software-based systems (pseudo-random numbers).

### True Random Number Generators (TRNGs)

TRNGs rely on physical processes, such as electronic noise or radioactive decay, to generate numbers that are truly random. These are less common in everyday computing due to the need for specialized hardware.

#### Advantages:

- Truly random, not predictable.
- Useful for cryptography and security.

### Disadvantages:

- Slower and less accessible.
- Requires specific hardware.

## Pseudo-Random Number Generators (PRNGs)

PRNGs are algorithms that use deterministic processes to generate sequences of numbers that appear random. These sequences are reproducible if the initial seed value is known.

### Advantages:

- Fast and easily implemented in software.
- Reproducible, which is useful for simulations and debugging.

### Disadvantages:

- Not truly random; sequences can eventually repeat.
- Predictable if the algorithm and seed are known, making it unsuitable for cryptographic purposes.

## Common PRNG Algorithms

### 1. Linear Congruential Generator (LCG):

- One of the oldest and simplest methods.
- Generates numbers using the formula:

$$X_{n+1} = (aX_n + c) \mod m$$

- Parameters  $a$ ,  $c$ , and  $m$  determine the quality and period of the sequence.

### 2. Mersenne Twister:

- Widely used due to its long period (e.g.,  $2^{19937} - 1$ ) and good statistical properties.
- Default PRNG in many programming environments, including Python.

### 3. Cryptographically Secure PRNGs (CSPRNGs):

- Designed to be unpredictable and secure, even if part of the sequence or the seed is known.
- Commonly used in cryptographic applications.

## Limitations in Software like Excel

Excel uses a PRNG to generate random numbers, specifically based on an LCG or variants, which has certain limitations:

### 1. Quality of Randomness:

- Excel's PRNG might not pass all statistical tests for randomness, especially in large-scale simulations.
- LCGs, particularly if not well-parameterized, can exhibit patterns or correlations.

### 2. Period Length:

- The period (the number of random numbers generated before the sequence repeats) can be limited in Excel, leading to potential issues in applications requiring large numbers of random values.

### 3. Reproducibility:

- While reproducibility is an advantage in simulations, it becomes a limitation in situations where true randomness is required, such as cryptography.

### 4. Limited Seed Control:

- Excel's PRNG allows limited control over the seed value, which can restrict reproducibility and customization.

### 5. Correlations:

- PRNGs like those in Excel can sometimes produce sequences with subtle correlations, which can affect the outcome of simulations, particularly in finance and risk modeling.

### 6. Lack of Advanced Features:

- Excel lacks advanced random number generation features available in specialized software, such as generating random numbers from various distributions directly, controlling for correlations, or using cryptographic PRNGs.

For more robust random number generation, specialized tools or programming environments like Python (using libraries such as `numpy`, `random`, or `secrets`) or dedicated statistical software (e.g., R, MATLAB) are often preferred.

## How to generate random numbers from any distribution

**Inverse transform sampling** is a general technique used to generate random samples from any probability distribution using its cumulative distribution function (CDF).

For distributions such as the **standard normal distribution**, the CDF does not have a closed-form expression, but the method works conceptually as follows:

#### 1. Generate a Uniform Random Variable $U$ :

Generate a random number  $U$  uniformly distributed in the interval  $[0, 1]$ . Those numbers are easy to generate as many programming languages and libraries have built in functions for this:

- **Python:** `np.random.uniform(0, 1)`
- **R:** `runif(1, min = 0, max = 1)`
- **Excel:** `RAND()`
- **Matlab:** `rand()`

#### 2. Apply the Inverse of the Standard Normal CDF:

Transform  $U$  using the inverse of the standard normal CDF. For example for the normal distribution  $Z = \Phi^{-1}(U)$  is a standard normal random variable.

## Normal distribution (Gaussian distribution)

To generate random numbers from a normal distribution with mean  $\mu$  and standard deviation  $\sigma$ , there are several methods.



## Box-Muller Transform

The **Box-Muller transform** is a widely used method to generate two independent standard normal random variables from two independent uniform random variables.

### Procedure:

1. **Generate Two Independent Uniform Random Variables  $U_1$  and  $U_2$ :**

These are uniformly distributed over  $(0, 1]$ .

2. **Transform Using Box-Muller Equations:**

$$Z_1 = \sqrt{-2 \ln U_1} \cdot \cos(2\pi U_2)$$

$$Z_2 = \sqrt{-2 \ln U_1} \cdot \sin(2\pi U_2)$$

Here,  $Z_1$  and  $Z_2$  are two independent standard normal random variables.

### Advantages:

- Simple and efficient.
- Generates two random normals at once.

### Implementation Example in Python:

PYTHON

```
import numpy as np

# Generate uniform random numbers
U1 = np.random.rand()
U2 = np.random.rand()

# Box-Muller transform
Z1 = np.sqrt(-2 * np.log(U1)) * np.cos(2 * np.pi * U2)
Z2 = np.sqrt(-2 * np.log(U1)) * np.sin(2 * np.pi * U2)
```

## Marsaglia Polar Method

The **Marsaglia polar method** is an optimized version of the Box-Muller transform that avoids trigonometric function evaluations, making it faster.

### Procedure:

1. **Generate Two Independent Uniform Random Variables  $U_1$  and  $U_2$ :**

Generate random numbers  $U_1, U_2$  uniformly over  $[-1, 1]$ .

2. **Compute  $S = U_1^2 + U_2^2$ :**

- If  $S \geq 1$  or  $S = 0$ , discard  $U_1$  and  $U_2$  and generate new values.
- If  $0 < S < 1$ , proceed.

3. **Transform:**

$$Z_1 = U_1 \sqrt{\frac{-2 \ln S}{S}}$$

$$Z_2 = U_2 \sqrt{\frac{-2 \ln S}{S}}$$

$Z_1$  and  $Z_2$  are independent standard normal random variables.

### Implementation Example in Python:

PYTHON

```
import numpy as np

while True:
    U1 = np.random.uniform(-1, 1)
    U2 = np.random.uniform(-1, 1)
    S = U1**2 + U2**2
    if S >= 1 or S == 0:
        continue
    break

factor = np.sqrt(-2 * np.log(S) / S)
Z1 = U1 * factor
Z2 = U2 * factor
```

### Ziggurat Algorithm

The **Ziggurat algorithm** is a more advanced method for generating standard normal random variables. It divides the probability density function (PDF) of the normal distribution into layered "strips" (resembling a ziggurat) and uses a combination of uniform sampling and rejection sampling to generate normal variates.

#### Advantages:

- Very efficient, especially for large-scale simulations.
- Faster than methods relying on trigonometric functions.

#### Disadvantages:

- More complex to implement.
- Requires precomputation of tables or setup.

Many statistical libraries implement this method due to its efficiency, for example, `numpy` uses the Ziggurat method for generating normal random variables by default in its `np.random.normal()` function.

### Using Built-in Libraries

Most programming languages and statistical libraries provide built-in functions to generate standard normal random variables. These functions often use optimized algorithms like the Ziggurat method or well-tested combinations of other methods for efficiency and speed.

## MATLAB:

MATLAB

```
% Generate a single random number with mean mu and standard deviation sigma
mu = 0;
sigma = 1;
x = mu + sigma * randn();

% Generate an array of random numbers
x_array = mu + sigma * randn(1, 1000); % 1000 random numbers
```

## Python:

PYTHON

```
import numpy as np
x_array = np.random.normal(loc=mu, scale=sigma, size=1000)
```

## R:

R

```
x_array <- rnorm(1000, mean = mu, sd = sigma)
```

## Excel:

```
=NORM.INV(RAND(), mean, standard_dev)
```

## Exponential Distribution

To generate random numbers from an exponential distribution with rate parameter  $\lambda$ :

### 1. Use the Inverse Transform Sampling:

If  $U$  is a uniform random variable in  $[0, 1]$ :

$$X = -\frac{1}{\lambda} \ln(1 - U)$$

Since  $1 - U$  is also uniformly distributed, this simplifies to:

$$X = -\frac{1}{\lambda} \ln(U)$$

## MATLAB:

MATLAB

```
% Generate a single random number with rate parameter lambda
lambda = 1;
x = -log(rand()) / lambda;

% Generate an array of random numbers
x_array = -log(rand(1, 1000)) / lambda; % 1000 random numbers
```

**Python:**

PYTHON

```
x_array = np.random.exponential(scale=1/lambda, size=1000)
```

**R:**

R

```
x_array <- rexp(1000, rate=lambda)
```

**Excel:**

```
=-LN(RAND()) / lambda
```

**Binomial Distribution**

To generate random numbers from a binomial distribution with number of trials  $n$  and probability  $p$ :

**MATLAB:**

MATLAB

```
% Generate a single random number with n trials and probability p
n = 10;
p = 0.5;
x = binornd(n, p);

% Generate an array of random numbers
x_array = binornd(n, p, 1, 1000); % 1000 random numbers
```

**Python:**

PYTHON

```
x_array = np.random.binomial(n, p, size=1000)
```

**R:**

R

```
x_array <- rbinom(1000, size=n, prob=p)
```

**Excel:**

```
=BINOM.INV(trials, probability, RAND())
```

## Poisson Distribution

To generate random numbers from a Poisson distribution with rate parameter  $\lambda$ :

**MATLAB:**

MATLAB

```
% Generate a single random number with rate parameter lambda
lambda = 5;
x = poissrnd(lambda);

% Generate an array of random numbers
x_array = poissrnd(lambda, 1, 1000); % 1000 random numbers
```

**Python:**

PYTHON

```
x_array = np.random.poisson(lam=lambda, size=1000)
```

**R:**

R

```
x_array <- rpois(1000, lambda=lambda)
```

**Excel:**

```
=POISSON.DIST(x, lambda, FALSE)
```

## Gamma Distribution

To generate random numbers from a Gamma distribution with shape parameter  $k$  and scale parameter  $\theta$ :

**MATLAB:**

MATLAB

```
% Generate a single random number with shape parameter k and scale
parameter theta
k = 2;
theta = 1;
x = gamrnd(k, theta);

% Generate an array of random numbers
x_array = gamrnd(k, theta, 1, 1000); % 1000 random numbers
```

**Python:**

PYTHON

```
x_array = np.random.gamma(shape=k, scale=theta, size=1000)
```

**R:**

R

```
x_array <- rgamma(1000, shape=k, scale=theta)
```

**Excel:**

```
=GAMMA.INV(RAND(), shape, scale)
```

**Beta Distribution**

To generate random numbers from a Beta distribution with shape parameters  $\alpha$  and  $\beta$ :

**MATLAB:**

MATLAB

```
% Generate a single random number with shape parameters alpha and beta
alpha = 2;
beta = 5;
x = betarnd(alpha, beta);

% Generate an array of random numbers
x_array = betarnd(alpha, beta, 1, 1000); % 1000 random numbers
```

**Python:**

```
x_array = np.random.beta(a=alpha, b=beta, size=1000)
```

**R:**

```
x_array %>% rbeta(1000, shape1=alpha, shape2=beta)
```

**Excel:**

```
=BETA.INV(RAND(), alpha, beta)
```

Dataview (for inline query '=NORM.INV(RAND(), mean, standard\_dev)'): Unrecognized function name 'RAND'

Dataview (inline field '-LN(RAND()) / lambda'): Error:

-- PARSING FAILED -----

```
> 1 | -LN(RAND()) / lambda
    | ^
```

Expected one of the following:

'(', 'null', boolean, date, duration, file link, list ('[1, 2, 3]'), negated field, number, object ('{ a: 1, b: 2 }'), string, variable

Dataview (for inline query '=BINOM.INV(trials, probability, RAND())'): Unrecognized function name 'RAND'

Dataview (for inline query '=POISSON.DIST(x, lambda, FALSE)'): Cannot call type 'null' as a function

Dataview (for inline query '=GAMMA.INV(RAND(), shape, scale)'): Unrecognized function name 'RAND'

Dataview (for inline query '=BETA.INV(RAND(), alpha, beta)'): Unrecognized function name 'RAND'

## Correlated standard normal random numbers

If we have two independent stochastic variables  $Z_1$  and  $Z_2$  that are both standard normally distributed. If we wish to transform them into two stochastic variables  $W_1$  and  $W_2$  that have a correlation coefficient  $\rho$ , it is clear how this can be done. Define:

$$\begin{cases} W_1 = Z_1 \\ W_2 = \rho Z_1 + \sqrt{1 - \rho^2} Z_2 \end{cases}$$

It is clear that the mean for both  $W_1$  and  $W_2$  is zero. Due to the independence of  $Z_1$  and  $Z_2$ , the variance of

$$\text{Var}(W_2) = \rho^2 \text{Var}(Z_1) + (1 - \rho^2) \text{Var}(Z_2) = 1$$

and

$$\text{Covar}(W_1, W_2) = \rho^2$$

To apply this for more than 2 random variables, we need to extend this idea. The solution is found into a particular decomposition of the correlation matrix: the Cholesky decomposition.

Let's go through a step-by-step proof that the Cholesky decomposition provides the correct correlation structure when generating correlated random variables.

**Objective:** Given  $n$  independent standard normal random variables  $Z_1, Z_2, \dots, Z_n$  and a desired correlation matrix  $\mathbf{C}$ , we want to create a new set of correlated normal variables  $W_1, W_2, \dots, W_n$  such that they have the desired correlation matrix.

Given a symmetric positive definite correlation matrix  $\mathbf{C}$ , there exists a unique lower triangular matrix  $\mathbf{L}$  such that:

$$\mathbf{C} = \mathbf{L}\mathbf{L}^\top$$

Here,  $\mathbf{L}$  is the Cholesky factor of  $\mathbf{C}$ .

Let  $\mathbf{Z}$  be a column vector of independent standard normal random variables:

$$\mathbf{Z} = \begin{pmatrix} Z_1 \\ Z_2 \\ \vdots \\ Z_n \end{pmatrix}$$

Define a new vector  $\mathbf{W}$  as:

$$\mathbf{W} = \mathbf{L}\mathbf{Z}$$

Where  $\mathbf{W}$  will be a vector of correlated normal random variables.

The covariance matrix  $\Sigma_{\mathbf{W}}$  of the vector  $\mathbf{W}$  is given by:

$$\Sigma_{\mathbf{W}} = \text{Cov}(\mathbf{W}) = \text{Cov}(\mathbf{L}\mathbf{Z})$$

Since  $\mathbf{L}$  is a deterministic matrix (non-random), we can express the covariance as:

$$\Sigma_{\mathbf{W}} = \mathbf{L} \cdot \text{Cov}(\mathbf{Z}) \cdot \mathbf{L}^\top$$

Since the components of  $\mathbf{Z}$  are independent standard normal random variables, the covariance matrix of  $\mathbf{Z}$  is the identity matrix  $\mathbf{I}$ :



$$\text{Cov}(\mathbf{Z}) = \mathbf{I}$$

Therefore:

$$\Sigma_{\mathbf{W}} = \mathbf{L}\mathbf{L}\mathbf{L}^{\top} = \mathbf{L}\mathbf{L}^{\top}$$

Recall that  $\mathbf{C} = \mathbf{L}\mathbf{L}^{\top}$ , so:

$$\Sigma_{\mathbf{W}} = \mathbf{C}$$

This shows that the covariance matrix of  $\mathbf{W}$  (the correlated variables) is exactly  $\mathbf{C}$ , the desired correlation matrix.

Therefore, by applying the Cholesky decomposition to the correlation matrix and using it to transform independent standard normal variables, you obtain correlated normal variables with the correct correlation structure. This gives a technique to generate correlated random variables.

## Discretization of non-Brownian processes

### Discretization of non-Brownian processes

#### Jump Diffusion

Discretizing a **jump diffusion model** involves approximating both the continuous Brownian motion component and the discontinuous jump component of an SDE. For example, the classic jump diffusion model, proposed by Robert C. Merton, extends the GBM to include jumps:

$$dS_t = \mu S_t dt + \sigma S_t dW_t + S_t dJ_t$$

where:

- $S_t$  is the asset price at time  $t$ .
- $\mu$  is the drift rate.
- $\sigma$  is the volatility of the continuous part of the asset's return.
- $W_t$  is a standard Brownian motion.
- $J_t$  is a jump process, typically modeled as a **Poisson process** with intensity  $\lambda$  (jump frequency) and jump size  $Y$  (typically log-normally distributed).

The term  $S_t dJ_t$  represents the jumps in the asset price, with:

- **Jump Intensity**  $\lambda$ : The average number of jumps per unit time.
- **Jump Size**  $Y$ : The multiplicative effect of a jump on the asset price, often  $Y = e^Z - 1$  where  $Z \sim N(\nu, \delta^2)$ .

The continuous part, represented by the Brownian motion term, is discretized in a manner similar to GBM. Using the Euler method for the continuous component:

$$S_{t+\Delta t} = S_t \left( 1 + \mu\Delta t + \sigma\sqrt{\Delta t} Z \right)$$

or in a multiplicative form (more typical for financial modeling):

$$S_{t+\Delta t}^{\text{cont}} = S_t \exp \left( \left( \mu - \frac{1}{2}\sigma^2 \right) \Delta t + \sigma\sqrt{\Delta t} Z \right)$$

where  $Z \sim N(0, 1)$  is a standard normal random variable.

For the jump component, we need to:

- Determine whether a jump occurs in each time step  $\Delta t$ .
- If a jump occurs, determine the size of the jump.

The occurrence of jumps is modeled by a Poisson process. The probability of one or more jumps occurring in a small time interval  $\Delta t$  is given by  $\lambda \Delta t$ . To determine if a jump occurs:

1. **Generate a Uniform Random Variable**  $U \sim \text{Uniform}(0, 1)$ .
2. **Check if**  $U < \lambda \Delta t$ :
  - If true, a jump occurs in this time step.
  - If false, no jump occurs.

If a jump occurs, the size of the jump  $Y$  is drawn from the specified distribution. Typically, the log of the jump size  $\ln(1 + Y) = Z$  is normally distributed:

$$Z \sim N(\nu, \delta^2)$$

Thus, the jump size  $Y = e^Z - 1$ .

Once we have both the continuous and jump components, we combine them to update the asset price. For each time step:

1. **Update the asset price with the continuous component:**

$$S_{t+\Delta t}^{\text{cont}} = S_t \exp \left( \left( \mu - \frac{1}{2} \sigma^2 \right) \Delta t + \sigma \sqrt{\Delta t} Z \right)$$

2. **Check for a jump and update accordingly:**

- If no jump occurs:

$$S_{t+\Delta t} = S_{t+\Delta t}^{\text{cont}}$$

- If a jump occurs:

$$S_{t+\Delta t} = S_{t+\Delta t}^{\text{cont}} \times (1 + Y)$$

where  $Y = e^Z - 1$ , and  $Z \sim N(\nu, \delta^2)$ .

This final step accounts for the combined effect of both the continuous and jump components in the discretized asset price.

## Heston model

Heston and other stochastic volatility models are governed by 2 SDEs: one for the price process and one for the variance process. Both need to be discretized together, taking care of the correlation as well.

Discretizing the Heston Stochastic Differential Equation (SDE) is an important step in simulating the dynamics of financial assets with stochastic volatility. The Heston model is widely used in finance because it captures the empirical observation that volatility is not constant but varies over time, often in a mean-reverting manner.

The Heston model describes the dynamics of the asset price  $S_t$  and its variance  $v_t$  using the following SDEs:

$$\begin{aligned} dS_t &= \mu S_t dt + \sqrt{v_t} S_t dW_t^S, \\ dv_t &= \kappa(\eta - v_t) dt + \theta \sqrt{v_t} dW_t^v, \end{aligned}$$

where:

- $S_t$  is the asset price at time  $t$ .
- $v_t$  is the variance of the asset at time  $t$  (volatility squared).
- $\mu$  is the drift rate of the asset price.
- $\kappa$  is the rate of mean reversion of the variance.
- $\eta$  is the long-term mean level of the variance.
- $\theta$  is the volatility of the variance (vol of vol).
- $W_t^S$  and  $W_t^v$  are two Wiener processes (Brownian motions) with correlation  $\rho$ .

There are several methods for discretizing the Heston SDE, with the most common being the **Euler method** and the **Milstein method**. Due to the nature of the variance process (which must remain positive), certain techniques are often employed to ensure numerical stability and accuracy.

The Euler method is the simplest discretization technique, though it may require some adjustments to handle the variance process effectively.

For a time step  $\Delta t$ , the discretized Heston SDEs are:

$$\begin{aligned} S_{t+\Delta t} &= S_t + \mu S_t \Delta t + \sqrt{v_t} S_t \sqrt{\Delta t} Z_1, \\ v_{t+\Delta t} &= v_t + \kappa(\eta - v_t) \Delta t + \theta \sqrt{v_t} \sqrt{\Delta t} Z_2, \end{aligned}$$

where:

- $Z_1$  and  $Z_2$  are standard normal random variables with correlation  $\text{Corr}(Z_1, Z_2) = \rho$ .

The variance process  $v_t$  must be positive, so special care is needed to avoid negative values.

### Handling the Variance Process:

Due to the discretization of the processes, it is possible for the variance to become negative in a sample path. There are multiple fixes for this problem. Note that any of such fixes, alters the underlying dynamic distribution ever so slightly. The smaller the time step is, the less likely this problem occurs. So if a lot of sample paths suffer from this problem, the time step should be reduced first.

But since this problem is inevitable, the following practical solutions have been proposed"

- **Reflection:** Alternatively, the variance can be reflected back into the positive domain if it becomes negative.
- **Absorption** if the variance becomes negative, just set it back to zero.
- **Truncation:** If  $v_{t+\Delta t}$  becomes negative, it can be truncated as follows:

$$\begin{aligned} v_{t_i+1}^* &= v_{t_i}^* + \kappa(\eta - v_{t_i})\Delta t + \lambda \sqrt{v_{t_i}} \sqrt{\Delta t} \varepsilon_2 \\ v_{t_i+1} &= \max(0, v_{t_i+1}^*) \end{aligned}$$

- **Partial Truncation:** we do not need to alter the entire variance, just the part that creates the problem:

$$\begin{aligned}
v_{t_{i+1}}^* &= v_{t_i}^* + \kappa(\eta - v_{t_i}^*)\Delta t + \lambda\sqrt{v_{t_i}}\sqrt{\Delta t}\varepsilon_2 \\
v_{t_{i+1}} &= \max(0, v_{t_{i+1}}^*)
\end{aligned}$$

The Milstein method improves upon Euler by including an additional term that captures the derivative of the diffusion function, reducing discretization error.

For the Heston model, the Milstein discretization for the variance process  $v_t$  is:

$$v_{t+\Delta t} = v_t + \kappa(\eta - v_t)\Delta t + \theta\sqrt{v_t}\sqrt{\Delta t}Z_2 + \frac{\theta^2\Delta t}{4}(Z_2^2 - 1)$$

The asset price  $S_t$  is discretized similarly to the Euler method.

We can summarize all this into the full Discretization Procedure:

#### 1. **Generate Correlated Brownian Motions:**

Generate two standard normal random variables  $Z_1$  and  $Z_2$  with correlation  $\rho$  using:

$$Z_2 = \rho Z_1 + \sqrt{1 - \rho^2} Z_3$$

where  $Z_3$  is another independent standard normal variable.

#### 2. **Discretize the Variance Process:**

Using either the Euler or Milstein method, update the variance  $v_{t+\Delta t}$ .

For Euler:

$$v_{t+\Delta t} = v_t + \kappa(\theta - v_t)\Delta t + \theta\sqrt{v_t}\sqrt{\Delta t}Z_2$$

For Milstein:

$$v_{t+\Delta t} = v_t + \kappa(\theta - v_t)\Delta t + \theta\sqrt{v_t}\sqrt{\Delta t}Z_2 + \frac{\theta^2\Delta t}{4}(Z_2^2 - 1)$$

**Check for Positivity:** Ensure  $v_{t+\Delta t} \geq 0$  using truncation or reflection if necessary.

#### 3. **Discretize the Asset Price Process:**

Update the asset price  $S_{t+\Delta t}$ :

$$S_{t+\Delta t} = S_t + \mu S_t \Delta t + \sqrt{v_t} S_t \sqrt{\Delta t} Z_1$$

#### 4. **Iterate:**

Repeat the above steps for all time steps to generate the full path of  $S_t$  and  $v_t$ .

### Variance gamma model

For the **Variance Gamma (VG) process** there is no diffusion component ( $\sigma = 0$ ), and its jumps are governed by a gamma distribution. The characteristic function of the VG process is:

To discretize the VG process:

1. **Gamma Subordinator:** Simulate the gamma subordinator  $G_t$  with mean rate  $\nu^{-1}$  and variance rate  $\nu^{-1}$ . This determines the arrival times of jumps.
2. **Jump Sizes:** The sizes of jumps  $Y_i$  are normally distributed with mean  $\theta\nu$  and variance  $\sigma^2\nu$ .
3. **Update Process:** For each time step, update the process using:

$$X_{t+\Delta t} = X_t + \sum_{i=1}^{N_{\Delta t}} Y_i$$

where  $N_{\Delta t}$  is the number of jumps in  $\Delta t$  from the gamma process.

## Multivariate discretization

To be written

## The Brownian bridge

A **Brownian bridge** is a stochastic process that is derived from a standard Brownian motion (or Wiener process) but with an additional condition that it must return to a specific value (usually zero) at a fixed future time. It can be thought of as a Brownian motion that is "pinned" at both the starting time and a fixed endpoint. This makes the Brownian bridge a useful tool in various applications, including financial mathematics, statistics, and physics.

While Brownian motion is unconstrained and can take any value over time, the Brownian bridge is constrained to return to a specific value (typically zero) at a fixed time  $T$ . This makes the Brownian bridge useful for modeling scenarios where an initial condition and a terminal condition are known or required.

Although we had already discussed an example of a constrained Brownian motion in the [OU Model](#), the Brownian bridge is pinned to specific values at a fixed time, rather than just have the weaker mean reversion constraint from the OU model.

### Definition

Formally, if  $W(t)$  is a standard Brownian motion starting at 0, then the Brownian bridge  $X(t)$  over the interval  $[0, T]$  is defined as:

$$X(t) = W(t) - \frac{t}{T}W(T), \quad 0 \leq t \leq T$$

where:

- $W(t)$  is a Brownian motion at time  $t$ ,
- $T$  is the fixed endpoint time.

### Properties

1. **Mean:** The mean of the Brownian bridge is zero for all  $t$ :

$$\mathbb{E}[X(t)] = 0$$

2. **Covariance:** The covariance between two points  $t_1$  and  $t_2$  in the Brownian bridge is given by:

$$\text{Cov}(X(t_1), X(t_2)) = t_1 \left(1 - \frac{t_2}{T}\right), \quad \text{for } 0 \leq t_1 \leq t_2 \leq T$$

3. **Variance:** The variance at time  $t$  is:

$$\text{Var}(X(t)) = t \left(1 - \frac{t}{T}\right)$$

This indicates that the variance starts at 0 at  $t = 0$ , increases to a maximum at  $t = T/2$ , and then decreases back to 0 at  $t = T$ .

The Brownian bridge is also related to other processes, such as the Ornstein-Uhlenbeck process, in that it represents a type of constrained Brownian motion.

Certainly! Below is a MATLAB code snippet to generate sample paths of a Brownian bridge, where you can specify different start and end values.

MATLAB

```
% Parameters
T = 1;           % Time horizon
N = 1000;        % Number of time steps
dt = T/N;        % Time step size
t = linspace(0, T, N+1); % Time vector

startValue = 5;   % Starting value of the bridge
endValue = 2;     % Ending value of the bridge

% Preallocate Brownian Bridge array
X = zeros(1, N+1);

% Generate a Brownian motion path
W = [0, cumsum(sqrt(dt) * randn(1, N))];

% Construct the Brownian Bridge with specified start and end values
for i = 1:N+1
    X(i) = startValue + W(i) - (t(i)/T) * (W(end) - endValue + startValue);
end

% Plot the sample path
figure;
plot(t, X, 'b-', 'LineWidth', 1.5);
hold on;
plot([0 T], [startValue endValue], 'ro'); % Mark the start and end points
xlabel('Time');
ylabel('Value');
title('Sample Path of a Brownian Bridge');
grid on;
```

## How It Works:

- **T**: The total time horizon of the Brownian bridge.
- **N**: The number of time steps in the discretization.
- **startValue** and **endValue**: The initial and final values of the Brownian bridge, which can be different.
- **W**: A standard Brownian motion path generated using cumulative sums of normal random variables scaled by **sqrt(dt)**.
- The Brownian bridge **X** is constructed by adjusting the Brownian motion path **W** so that it starts at **startValue** and ends at **endValue**.

You can modify the **startValue** and **endValue** to generate different Brownian bridges. The plot will show the sample path of the Brownian bridge along with markers indicating the start and end points.

## Applications

The Brownian bridge has various applications in finance, statistics, and physics. In finance, it is particularly useful in modeling and pricing derivative securities, such as barrier options, and in risk management. Here are some detailed applications:

### Barrier Options Pricing

Barrier options are path-dependent options where the payoff depends on whether the underlying asset's price crosses a certain level (the barrier) during the option's life. There are two main types of barrier options:

- **Knock-in options**: Become active only if the underlying asset crosses the barrier.
- **Knock-out options**: Become void if the underlying asset crosses the barrier.

### Application of Brownian Bridge in Barrier Options:

- **Conditioned Simulation**: When simulating the underlying asset's path for barrier options, the Brownian bridge can be used to conditionally simulate the asset's path so that it stays within certain bounds. For instance, when generating a path for a knock-out option, if you already know the start and end points of the asset's price, you can use a Brownian bridge to generate the intermediate prices while ensuring the barrier isn't breached.
- **Probability Estimation**: The Brownian bridge can be used to estimate the probability that the underlying asset will hit the barrier. Given the start and end prices of the underlying asset, the Brownian bridge allows for more efficient calculation of the probability that the path will hit the barrier, which is crucial in pricing barrier options.

#### Example:

For a down-and-out call option (a type of knock-out option where the option is void if the price goes below a certain barrier  $L$ ), you can estimate the probability that the price touches the barrier before maturity using the Brownian bridge:

$$\text{Prob}(\min_{0 \leq t \leq T} S(t) \leq L) \approx \text{Prob}(S(T) \leq L) + \frac{S(0) - L}{S(0) - S(T)} \text{Prob}(S(T) > L)$$

This probability can then be used to adjust the standard Black-Scholes option pricing formula.

## Lookback Options Pricing

Lookback options are another type of path-dependent option, where the payoff depends on the maximum or minimum price of the underlying asset over the option's life.

### Application of Brownian Bridge in Lookback Options:

- The Brownian bridge is used to estimate the distribution of the maximum or minimum of the underlying asset's price over a period. Since the Brownian bridge allows for conditioning on the start and end points of the asset's path, it can be used to simulate the entire path while keeping track of the maximum or minimum value.
- This is particularly useful in pricing floating-strike lookback options, where the strike price is determined as the maximum or minimum price of the underlying asset during the option's life.

## Asian Options Pricing

Asian options are options where the payoff depends on the average price of the underlying asset over a certain period. These options are less sensitive to short-term volatility spikes than standard options.

### Application of Brownian Bridge in Asian Options:

- In pricing Asian options, especially when using Monte Carlo simulations, the Brownian bridge is used to efficiently simulate the underlying asset's path, given certain boundary conditions. The bridge ensures that the simulated paths are consistent with the known start and end prices, which can improve the accuracy and efficiency of the simulation.

## Goodness-of-Fit Tests

In statistics, the Brownian bridge is the limiting distribution of certain test statistics used in goodness-of-fit tests, such as the Kolmogorov-Smirnov test. This test assesses whether a sample comes from a specified distribution.

### Application of Brownian Bridge in Goodness-of-Fit Tests:

- When conducting a Kolmogorov-Smirnov test, the test statistic (which measures the maximum deviation between the empirical distribution function of the sample and the cumulative distribution function of the reference distribution) converges in distribution to a Brownian bridge under the null hypothesis.
- The Brownian bridge is thus used to derive the distribution of the test statistic under the null hypothesis, allowing for the calculation of p-values and making decisions about the fit of the data.

## Stochastic Modeling in Physics and Engineering

In physics and engineering, Brownian bridges model various phenomena where a system is conditioned to return to a specific state after a certain period. Examples include particle diffusion with absorbing boundaries and the modeling of fluctuations in physical systems constrained by boundary conditions.



## Application in Diffusion Processes:

- **Constrained Diffusion:** Brownian bridges model diffusion processes where a particle is required to return to a specific point after a certain time, such as in models of particle reabsorption or restricted diffusion in porous media.
- **Path Integral Methods:** In quantum mechanics, Brownian bridges are used in path integral formulations where the paths of particles are constrained to start and end at specific points, relevant in fields like quantum field theory and statistical mechanics.

## Variance reduction methods

Variance reduction methods are techniques used in Monte Carlo simulations to increase the accuracy of the estimated results without increasing the number of simulations. These methods aim to reduce the variance of the estimator, leading to more precise estimates with fewer simulation runs.

Two excellent books that cover Monte carlo in a lot of detail are

- [Paul Glasserman - Monte Carlo Methods in Financial Engineering](#)
- [Peter Jäckel - Monte Carlo Methods in Finance](#)

Some of the common variance reduction techniques are listed below and explained into more detail further. Each of these variance reduction techniques is aimed at reducing the variability in Monte Carlo estimates, thereby increasing their efficiency. These methods are particularly valuable in financial engineering when calculating the prices of complex derivatives or to evaluate risk metrics such as Value at Risk (VaR). By applying these methods, you can achieve more accurate results with fewer simulations, reducing computational cost and time. Obviously, some of those methods can be used together.

- **Antithetic Variates**

Antithetic variates involve generating pairs of dependent random variables in such a way that they are negatively correlated. The idea is that if one simulation run overestimates the result, the paired simulation might underestimate it, balancing the overall result.

- **Moment Matching**

Moment matching involves adjusting the sample so that it matches the desired moments (mean, variance, skewness, etc.) of the underlying distribution more closely.

- **Control Variates**

Control variates use the known properties of a related variable (the control) to reduce the variance of the estimator. If the control variate is correlated with the variable of interest and its expected value is known, it can be used to adjust the estimator.

- **Importance Sampling**

Importance sampling changes the probability distribution from which samples are drawn to make rare events more common or to focus on more important parts of the distribution. The goal is to increase the efficiency of the estimator by concentrating on the areas that contribute most to the final result.

- **Stratified Sampling**

Stratified sampling involves dividing the domain of the input random variables into different "strata" or subdomains and then sampling from each stratum. The idea is to ensure that each part of the domain is adequately represented in the sample.

- **Latin Hypercube Sampling**

Latin Hypercube Sampling (LHS) is a method that ensures that each parameter is sampled across its entire range in a stratified manner. Unlike simple random sampling, LHS ensures that the full range of each input variable is sampled more uniformly.

- **Quasi-Monte Carlo Methods**

Quasi-Monte Carlo (QMC) methods replace random numbers with deterministic sequences (low-discrepancy sequences) that are designed to cover the space more uniformly than random samples.

## Antithetic Variates

### Antithetic variates

Antithetic variates are a powerful variance reduction technique used in Monte Carlo simulations. The main idea is to reduce the variance of the estimator by generating pairs of negatively correlated random variables, which helps to balance out the variability in the estimates.

Antithetic variates is a straightforward yet powerful method to reduce variance in Monte Carlo simulations. By generating negatively correlated pairs of random variables, it ensures that overestimations and underestimations balance out, resulting in a more accurate and efficient estimator.

### Detailed Explanation of Antithetic Variates

#### Basic Concept

When using Monte Carlo methods to estimate an expected value  $\mathbb{E}[f(X)]$ , where  $X$  is a random variable and  $f(X)$  is some function of interest, the accuracy of the estimate depends on the variance of the estimator. Antithetic variates aim to reduce this variance by introducing negative correlation between paired simulations.

#### How it Works

Consider a random variable  $X$  uniformly distributed on  $[0, 1]$  and suppose you want to estimate  $\mathbb{E}[f(X)]$  using Monte Carlo simulation. The standard approach is to draw  $n$  independent samples  $X_1, X_2, \dots, X_n$  from the distribution of  $X$ , compute the function values  $f(X_1), f(X_2), \dots, f(X_n)$ , and then average them:

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n f(X_i)$$

In the antithetic variates method, for each sample  $X_i$  from the uniform distribution, we also consider the "antithetic" sample  $1 - X_i$ . The idea is to compute the function value for both  $X_i$  and  $1 - X_i$  and then average these two values:

$$\hat{\mu}_i^{\text{antithetic}} = \frac{f(X_i) + f(1 - X_i)}{2}$$

The overall estimator is then the average of these antithetic averages:

$$\hat{\mu}_{\text{antithetic}} = \frac{1}{n} \sum_{i=1}^n \hat{\mu}_i^{\text{antithetic}} = \frac{1}{2n} \sum_{i=1}^n (f(X_i) + f(1 - X_i))$$

## Mathematical Justification

The key to understanding why antithetic variates work lies in the correlation between  $f(X_i)$  and  $f(1 - X_i)$ . Ideally, we want  $f(X_i)$  and  $f(1 - X_i)$  to be negatively correlated, meaning when one is above its mean, the other tends to be below its mean. This negative correlation reduces the overall variance of the estimator.

Let's denote:

$$\text{Var}(\hat{\mu}) = \frac{1}{n} \text{Var}(f(X))$$

for the standard Monte Carlo estimator and

$$\text{Var}(\hat{\mu}_{\text{antithetic}}) = \frac{1}{4n} \text{Var}(f(X) + f(1 - X))$$

for the antithetic variates estimator.

Using the variance formula for the sum of two random variables  $A$  and  $B$ :

$$\text{Var}(A + B) = \text{Var}(A) + \text{Var}(B) + 2\text{Cov}(A, B)$$

we get:

$$\text{Var}(f(X) + f(1 - X)) = 2\text{Var}(f(X)) + 2\text{Cov}(f(X), f(1 - X))$$

So the variance of the antithetic estimator becomes:

$$\text{Var}(\hat{\mu}_{\text{antithetic}}) = \frac{1}{2n} \text{Var}(f(X)) + \frac{1}{2n} \text{Cov}(f(X), f(1 - X))$$

Since  $\text{Cov}(f(X), f(1 - X))$  is negative, the variance of the antithetic variates estimator  $\text{Var}(\hat{\mu}_{\text{antithetic}})$  is generally smaller than the variance of the standard Monte Carlo estimator.

## Example: Estimating the Value of a European Call Option

Suppose you want to estimate the price of a European call option with payoff  $\max(S_T - K, 0)$  using Monte Carlo simulation, where  $S_T$  is the asset price at maturity  $T$  and  $K$  is the strike price.

1. **Simulate the paths:** For each sample  $Z_i$  from a standard normal distribution, simulate the asset price using:

$$S_T^{(i)} = S_0 \exp \left( \left( r - \frac{\sigma^2}{2} \right) T + \sigma \sqrt{T} Z_i \right)$$

where  $S_0$  is the initial price,  $r$  is the risk-free rate, and  $\sigma$  is the volatility.

2. **Antithetic variates:** For each  $Z_i$ , also simulate the asset price using  $-Z_i$  to get:

$$S_T^{(i, \text{antithetic})} = S_0 \exp \left( \left( r - \frac{\sigma^2}{2} \right) T + \sigma \sqrt{T} (-Z_i) \right)$$

3. **Compute the payoffs:** For each pair  $(Z_i, -Z_i)$ , calculate the option payoffs:

$$\text{Payoff}_i = \frac{1}{2} \left( \max(S_T^{(i)} - K, 0) + \max(S_T^{(i, \text{antithetic})} - K, 0) \right)$$

4. **Estimate the option price:** Average these payoffs and discount them back to the present value:

$$\hat{C}_{\text{antithetic}} = e^{-rT} \cdot \frac{1}{n} \sum_{i=1}^n \text{Payoff}_i$$

## Benefits and Limitations

- **Benefits:**
  - **Variance Reduction:** The primary benefit is a reduction in variance, which improves the accuracy of the Monte Carlo estimate.
  - **Computational Efficiency:** Since antithetic pairs are typically generated from the same random sample, this method requires little additional computational effort.
- **Limitations:**
  - **Function Dependence:** The effectiveness of antithetic variates depends on the function  $f(X)$ . If  $f(X)$  and  $f(1 - X)$  are not sufficiently negatively correlated, the variance reduction might be minimal or even non-existent.
  - **No Universal Solution:** Not all functions or problems benefit equally from antithetic variates. For some cases, other variance reduction techniques might be more effective.

## Moment Matching

### Moment Matching

Moment Matching is a variance reduction technique used in Monte Carlo simulations to ensure that the simulated sample moments (e.g., mean, variance) match the theoretical moments of the underlying distribution. By adjusting the generated samples so that their moments match the expected moments, Moment Matching can lead to more accurate and reliable estimates, particularly when the distribution of the output is well understood.

While Moment Matching is most commonly applied to the first two moments (mean and variance), it can be extended to higher-order moments, although this increases the complexity of the adjustment. Despite its simplicity, Moment Matching is widely used in applications where accurate moment estimation is crucial, such as in option pricing, risk management, and the simulation of stochastic processes.

### Detailed Explanation of Moment Matching

#### Basic Concept

In Monte Carlo simulations, random samples are drawn from a probability distribution to estimate certain quantities, such as expected values or probabilities. However, due to the randomness of the sampling process, the sample moments (e.g., mean, variance) may differ from the theoretical moments of the distribution. This discrepancy can lead to biased estimates and increased variance.

Moment Matching corrects this by adjusting the samples so that their moments (up to a certain order, typically the first and second moments) align with the theoretical moments of the distribution. This adjustment reduces the variance of the estimator and can lead to more accurate simulation results.

## How it Works

Let's consider the case where we want to match the first two moments, the mean and variance, which is the most common application of Moment Matching.

### 1. Generate Initial Samples:

- Draw  $n$  independent samples  $X_1, X_2, \dots, X_n$  from the target distribution  $F(x)$  with mean  $\mu$  and variance  $\sigma^2$ .
- Compute the sample mean  $\hat{\mu}$  and sample variance  $\hat{\sigma}^2$  of the drawn samples:

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n X_i$$
$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \hat{\mu})^2$$

### 2. Adjust the Samples:

- Adjust the samples to match the theoretical mean  $\mu$  and variance  $\sigma^2$ . This is done by transforming the initial samples as follows:

$$Y_i = \mu + \frac{\sigma}{\hat{\sigma}} (X_i - \hat{\mu})$$

- Here,  $Y_i$  are the adjusted samples that have the desired mean and variance. Specifically, the transformation rescales the deviations of the samples from their mean and then shifts them so that the mean matches  $\mu$ .

### 3. Compute the Estimate:

- Use the adjusted samples  $Y_1, Y_2, \dots, Y_n$  to compute the final estimate of the quantity of interest, such as an expected value:

$$\hat{\mu}_{\text{MM}} = \frac{1}{n} \sum_{i=1}^n g(Y_i)$$

- Here,  $g(Y_i)$  represents the function being evaluated at each adjusted sample.

## Variance Reduction

The key advantage of Moment Matching is the reduction in variance of the estimator. By ensuring that the sample moments match the theoretical moments, the estimator becomes more accurate and less prone to the random fluctuations that can occur in standard Monte Carlo simulations.

The variance reduction is particularly noticeable when the function  $g(X)$  being estimated is sensitive to deviations in the mean or variance. In such cases, ensuring that the sample moments align with the theoretical moments can lead to significant improvements in accuracy.

### Example: Estimating the Mean of a Distribution

Suppose we want to estimate the mean  $\mu$  of a distribution using Monte Carlo simulation. In standard Monte Carlo, we would draw  $n$  independent samples  $X_1, X_2, \dots, X_n$  and compute the sample mean:

$$\hat{\mu}_{MC} = \frac{1}{n} \sum_{i=1}^n X_i$$

However, due to random sampling, the sample mean  $\hat{\mu}$  may differ from the true mean  $\mu$ , introducing variance in the estimate.

Using Moment Matching, we adjust the samples so that the mean and variance match the theoretical values:

1. **Compute the Sample Mean and Variance:**

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n X_i, \quad \hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \hat{\mu})^2$$

2. **Adjust the Samples:**

$$Y_i = \mu + \frac{\sigma}{\hat{\sigma}}(X_i - \hat{\mu})$$

3. **Estimate the Mean:**

$$\hat{\mu}_{MM} = \frac{1}{n} \sum_{i=1}^n Y_i$$

In this case,  $\hat{\mu}_{MM} = \mu$  by construction, and the variance of the estimator is reduced because the samples have been adjusted to better represent the true distribution.

## Extension to Higher-Order Moments

Moment Matching can be extended to match higher-order moments, such as skewness (third moment) and kurtosis (fourth moment), depending on the needs of the simulation. For example, if the distribution is known to be skewed, matching the skewness can lead to further variance reduction.

However, adjusting for higher-order moments becomes more complex and may require nonlinear transformations of the samples. In practice, most applications focus on matching the first two moments (mean and variance) because this provides a good balance between complexity and effectiveness.

## Benefits and Limitations

- **Benefits:**

- **Variance Reduction:** By aligning the sample moments with the theoretical moments, Moment Matching can significantly reduce the variance of the Monte Carlo estimator, leading to more accurate and reliable estimates.
- **Simplicity:** For matching the first two moments, the adjustment is straightforward and easy to implement, involving simple linear transformations of the samples.
- **Flexibility:** Moment Matching can be combined with other variance reduction techniques, such as Antithetic Variates or Stratified Sampling, to further improve accuracy.

- **Limitations:**

- **Dependency on Distribution Knowledge:** Moment Matching requires knowledge of the theoretical moments of the distribution, which may not always be available

or easy to compute.

- **Limited to Finite Moments:** If the distribution has infinite moments (e.g., a heavy-tailed distribution), Moment Matching may not be applicable.
- **Nonlinearity for Higher Moments:** Extending Moment Matching to higher-order moments involves more complex transformations and may not always lead to significant improvements.

## Control Variates

### Control Variates

Control variates are another powerful variance reduction technique used in Monte Carlo simulations. The idea is to exploit the known expected value of a related variable (the control variate) to reduce the variance of the estimator for the quantity of interest.

Control variates are a highly effective variance reduction technique. By leveraging the known expected value of a related variable, they reduce the variance of the estimator, leading to more accurate and efficient simulations.

### Detailed Explanation of Control Variates

#### Basic Concept

In Monte Carlo simulations, we often aim to estimate the expected value of a random variable  $\mathbb{E}[f(X)]$ , where  $X$  is a random variable and  $f(X)$  is a function of interest. The accuracy of the estimate can be improved by introducing a control variate—a random variable  $Y$  that is correlated with  $f(X)$  and has a known expected value  $\mathbb{E}[Y]$ .

The control variate method works by adjusting the estimate of  $\mathbb{E}[f(X)]$  based on the deviation of  $Y$  from its known mean. This adjustment helps to reduce the variance of the estimator.

#### How it Works

Let  $Y$  be a random variable for which the expected value  $\mathbb{E}[Y]$  is known. We simulate both  $f(X)$  and  $Y$ , and then adjust the estimate of  $\mathbb{E}[f(X)]$  using the formula:

$$\hat{\mu}_{cv} = \frac{1}{n} \sum_{i=1}^n [f(X_i) - \beta(Y_i - \mathbb{E}[Y])]$$

where  $\beta$  is a coefficient that minimizes the variance of the adjusted estimator.

#### Deriving the Optimal $\beta$

The goal is to find the value of  $\beta$  that minimizes the variance of the estimator  $\hat{\mu}_{cv}$ . The variance of  $\hat{\mu}_{cv}$  is given by:

$$\text{Var}(\hat{\mu}_{cv}) = \frac{1}{n} \text{Var}(f(X) - \beta(Y - \mathbb{E}[Y]))$$

Expanding the variance term using the variance formula:

$$\text{Var}(A - B) = \text{Var}(A) + \beta^2 \text{Var}(B) - 2\beta \text{Cov}(A, B)$$

we get:

$$\text{Var}(\hat{\mu}_{cv}) = \frac{1}{n} (\text{Var}(f(X)) + \beta^2 \text{Var}(Y) - 2\beta \text{Cov}(f(X), Y))$$

To minimize this variance, we take the derivative with respect to  $\beta$  and set it to zero:

$$\frac{\partial}{\partial \beta} \text{Var}(\hat{\mu}_{cv}) = 2\beta \text{Var}(Y) - 2\text{Cov}(f(X), Y) = 0$$

Solving for  $\beta$ , we get the optimal coefficient:

$$\beta^* = \frac{\text{Cov}(f(X), Y)}{\text{Var}(Y)}$$

Substituting  $\beta^*$  back into the variance of  $\hat{\mu}_{cv}$  gives:

$$\text{Var}(\hat{\mu}_{cv}) = \frac{1}{n} \left( \text{Var}(f(X)) - \frac{\text{Cov}(f(X), Y)^2}{\text{Var}(Y)} \right)$$

This expression shows that the variance of the control variates estimator is reduced compared to the variance of the simple Monte Carlo estimator, provided that  $f(X)$  and  $Y$  are positively correlated.

### Example: Estimating the Value of a European Call Option

Consider the problem of estimating the price of a European call option with payoff  $\max(S_T - K, 0)$ , where  $S_T$  is the asset price at maturity  $T$ , and  $K$  is the strike price.

#### 1. Choosing a Control Variate:

- A common choice for a control variate in option pricing is the underlying asset itself,  $S_T$ , because we know its expected value under the risk-neutral measure:

$$\mathbb{E}[S_T] = S_0 e^{rT}$$

- Here,  $S_0$  is the initial asset price, and  $r$  is the risk-free rate.

#### 2. Simulate Paths:

- Generate  $n$  independent samples  $S_T^{(i)}$  of the asset price at maturity using the standard Monte Carlo approach.

#### 3. Calculate Payoffs and Control Variate:

- For each simulated path, calculate the option payoff  $C_i = \max(S_T^{(i)} - K, 0)$ .
- Also compute the control variate  $Y_i = S_T^{(i)}$ .

#### 4. Determine Optimal $\beta$ :

- Estimate  $\beta^*$  using sample estimates of the covariance  $\text{Cov}(C_i, Y_i)$  and variance  $\text{Var}(Y_i)$ :

$$\hat{\beta}^* = \frac{\frac{1}{n} \sum_{i=1}^n (C_i - \bar{C})(Y_i - \bar{Y})}{\frac{1}{n} \sum_{i=1}^n (Y_i - \bar{Y})^2}$$

where  $\bar{C}$  and  $\bar{Y}$  are the sample means of  $C_i$  and  $Y_i$ , respectively.

#### 5. Apply the Control Variate Adjustment:

- Adjust each payoff using the control variate:



$$C_i^{cv} = C_i - \hat{\beta}^*(Y_i - S_0 e^{rT})$$

- Average these adjusted payoffs to get the control variate estimator:

$$\hat{C}_{cv} = e^{-rT} \cdot \frac{1}{n} \sum_{i=1}^n C_i^{cv}$$

## Benefits and Limitations

- **Benefits:**
  - **Variance Reduction:** When the control variate is well-chosen (highly correlated with  $f(X)$ ), the reduction in variance can be substantial.
  - **Improved Accuracy:** The adjustment made by the control variate leads to more accurate estimates without needing to increase the number of simulations.
  - **Flexibility:** Control variates can be applied to a wide range of problems where an appropriate related variable with a known expected value can be identified.
- **Limitations:**
  - **Selection of Control Variate:** The effectiveness of this method depends heavily on choosing an appropriate control variate. If the control variate is poorly chosen (weakly correlated with  $f(X)$ ), the method may not reduce variance effectively.
  - **Additional Computation:** Estimating the optimal  $\beta$  and computing the adjusted estimator require additional computation, though this is often offset by the reduction in variance.

## Importance Sampling

### Importance Sampling

Importance Sampling is another powerful variance reduction technique in Monte Carlo simulations, particularly useful when dealing with rare events or when the region of interest in the probability distribution has low probability. The core idea is to change the probability distribution from which samples are drawn, so that more samples fall in the important regions of the space, and then reweight the samples to correct for this change.

By reweighting the samples appropriately, this method reduces the variance of the estimator, making it particularly useful for estimating probabilities of rare events or integrals over regions with low probability density. However, the success of Importance Sampling critically depends on choosing an appropriate importance distribution, which can be a challenging task.

### Detailed Explanation of Importance Sampling

#### Basic Concept

When estimating the expected value  $\mathbb{E}[f(X)]$  of a function  $f(X)$  under a probability distribution  $p(x)$ , the standard Monte Carlo estimator is given by:

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n f(X_i)$$

where  $X_i$  are independent samples drawn from the distribution  $p(x)$ .

In cases where  $f(X)$  is significant only in regions where  $p(x)$  is small, the variance of the estimator can be large because most of the samples  $X_i$  contribute little to the estimate. Importance Sampling addresses this issue by sampling from a different distribution  $q(x)$  that places more emphasis on the important regions, and then correcting for the change in distribution.

## How it Works

Let  $q(x)$  be a probability density function that is non-zero wherever  $p(x)f(x)$  is non-zero. The expected value  $\mathbb{E}[f(X)]$  under the original distribution  $p(x)$  can be rewritten as:

$$\mathbb{E}[f(X)] = \int f(x)p(x) dx = \int f(x)\frac{p(x)}{q(x)}q(x) dx = \mathbb{E}_q \left[ f(X) \frac{p(X)}{q(X)} \right]$$

where  $\mathbb{E}_q[\cdot]$  denotes the expectation under the new distribution  $q(x)$ .

The Importance Sampling estimator is then given by:

$$\hat{\mu}_{\text{IS}} = \frac{1}{n} \sum_{i=1}^n f(X_i) \frac{p(X_i)}{q(X_i)}$$

where  $X_i$  are samples drawn from  $q(x)$  instead of  $p(x)$ .

## Variance Reduction

The effectiveness of Importance Sampling lies in reducing the variance of the estimator. The variance of the Importance Sampling estimator is:

$$\text{Var}(\hat{\mu}_{\text{IS}}) = \frac{1}{n} \text{Var}_q \left( f(X) \frac{p(X)}{q(X)} \right)$$

which depends on the choice of  $q(x)$ . An ideal choice of  $q(x)$  minimizes this variance, making the estimator more efficient.

## Choosing the Importance Sampling Distribution $q(x)$

The key to effective Importance Sampling is choosing a good importance distribution  $q(x)$ . The optimal choice, which minimizes the variance, is:

$$q^*(x) \propto |f(x)|p(x)$$

This choice ensures that more samples are drawn from regions where the product  $|f(x)|p(x)$  is large, reducing the variance of the estimator. However, in practice, directly using  $q^*(x)$  is often difficult because it requires knowledge of  $f(x)$  and  $p(x)$  across the entire domain.

Therefore, in practice,  $q(x)$  is chosen as a distribution that is easier to sample from but still places more weight on the important regions where  $f(x)p(x)$  is large.

## Example: Estimating the Tail Probability of a Normal Distribution

Suppose we want to estimate the probability that a standard normal random variable  $Z$  exceeds a large threshold  $a$ , i.e.,  $\mathbb{P}(Z > a)$ . This is a rare event when  $a$  is large, so directly estimating it via Monte Carlo would require many samples, most of which would not contribute much to the estimate.

### 1. Standard Monte Carlo Estimation:

- Directly estimate the probability as:

$$\hat{P}_{MC} = \frac{1}{n} \sum_{i=1}^n \mathbf{1}_{\{Z_i > a\}}$$

where  $Z_i$  are i.i.d. samples from the standard normal distribution  $\mathcal{N}(0, 1)$ , and  $\mathbf{1}_{\{Z_i > a\}}$  is an indicator function that is 1 if  $Z_i > a$  and 0 otherwise.

- This estimator has high variance for large  $a$  because very few samples will satisfy  $Z_i > a$ .

## 2. Importance Sampling Approach:

- Choose a different distribution  $q(x)$ , such as a normal distribution with a shifted mean  $\mu = a$ , i.e.,  $q(x) = \mathcal{N}(a, 1)$ . This distribution places more weight on the tail where  $Z > a$ .
- Under this new distribution  $q(x)$ , estimate the probability as:

$$\hat{P}_{IS} = \frac{1}{n} \sum_{i=1}^n \mathbf{1}_{\{Z_i > a\}} \frac{p(Z_i)}{q(Z_i)}$$

where  $p(x)$  and  $q(x)$  are the probability density functions of the standard normal and shifted normal distributions, respectively. Specifically:

$$\frac{p(Z_i)}{q(Z_i)} = \exp\left(-\frac{1}{2}Z_i^2 + \frac{(Z_i - a)^2}{2}\right)$$

- Since samples  $Z_i$  are drawn from  $\mathcal{N}(a, 1)$ , a larger proportion of them will satisfy  $Z_i > a$ , leading to a lower-variance estimate.

## 3. Effectiveness:

- The variance of  $\hat{P}_{IS}$  is significantly reduced compared to  $\hat{P}_{MC}$ , making the estimation more accurate with fewer samples.

## Benefits and Limitations

### • Benefits:

- **Variance Reduction:** Properly chosen importance sampling distributions can lead to substantial variance reduction, especially in problems involving rare events.
- **Efficiency:** By focusing sampling efforts on important regions, Importance Sampling can achieve accurate results with fewer simulations, making it computationally efficient.

### • Limitations:

- **Choosing  $q(x)$ :** The effectiveness of Importance Sampling heavily depends on the choice of  $q(x)$ . If  $q(x)$  is poorly chosen, it can lead to increased variance or biased estimates.
- **Complexity:** Implementing Importance Sampling requires careful consideration of the problem and may involve more complex sampling procedures.

## Stratified Sampling

### Stratified Sampling

Stratified Sampling is a variance reduction technique used in Monte Carlo simulations to ensure that samples are spread more evenly across the entire range of possible values, rather than being clustered by chance in certain regions. This technique is particularly effective when the function being integrated or the region of interest has significant variability across different parts of the sample space.

The technique splits the sample space into strata and ensuring that each stratum is adequately sampled. This approach reduces the variance of the estimator, especially in cases where the function or distribution has significant variability across the space. While it adds some complexity to the simulation process, the benefits in terms of increased accuracy and reduced variance often outweigh these costs, making Stratified Sampling a valuable tool in many Monte Carlo applications.

### Detailed Explanation of Stratified Sampling

#### Basic Concept

In standard Monte Carlo simulation, samples are drawn independently from the entire distribution, which can lead to random clustering of samples in some areas and gaps in others. This randomness can cause the estimator to have high variance, especially when the function being integrated is not smooth or has significant variation in different regions.

Stratified Sampling addresses this issue by dividing the sample space into several non-overlapping "strata" or subregions and then sampling from each stratum. This approach ensures that each part of the sample space is represented in the final estimate, leading to a reduction in variance.

#### How it Works

##### 1. Partitioning the Sample Space:

- Divide the sample space into  $M$  non-overlapping strata  $S_1, S_2, \dots, S_M$ .
- The goal is to create strata such that within each stratum, the function being estimated has less variability. This often means dividing the space based on the cumulative distribution function (CDF) to ensure equal probability in each stratum.

##### 2. Sampling Within Each Stratum:

- For each stratum  $S_j$ , draw  $n_j$  samples independently from the distribution restricted to that stratum.
- The number of samples  $n_j$  in each stratum can be chosen based on the size of the stratum or other criteria, such as the variability of the function in that stratum.

##### 3. Combining the Results:

- Compute the estimate within each stratum  $S_j$ :

$$\hat{\mu}_j = \frac{1}{n_j} \sum_{i=1}^{n_j} f(X_{i,j})$$

where  $X_{i,j}$  are the samples drawn from stratum  $S_j$ .

- The overall estimate is then a weighted sum of the estimates from each stratum:

$$\hat{\mu}_{\text{Stratified}} = \sum_{j=1}^M w_j \hat{\mu}_j$$

where  $w_j = \mathbb{P}(X \in S_j)$  is the weight corresponding to the proportion of the total probability mass in stratum  $S_j$ .

## Variance Reduction

The variance of the stratified estimator can be expressed as:

$$\text{Var}(\hat{\mu}_{\text{Stratified}}) = \sum_{j=1}^M \frac{w_j^2}{n_j} \text{Var}(f(X) \mid X \in S_j)$$

This variance is generally lower than that of the standard Monte Carlo estimator because the samples are more evenly distributed across the entire sample space, reducing the chance of large deviations caused by random clustering.

## Optimal Allocation of Samples

In stratified sampling, an important consideration is how to allocate the total number of samples  $n = \sum_{j=1}^M n_j$  among the strata. The goal is to minimize the overall variance. Two common approaches are:

### 1. Proportional Allocation:

- Allocate samples in proportion to the size of the strata:

$$n_j = \frac{n \cdot w_j}{\sum_{k=1}^M w_k}$$

- This is simple and effective when the variability within each stratum is similar.

### 2. Optimal Allocation (Neyman Allocation):

- Allocate samples based on the variability within each stratum:

$$n_j = \frac{n \cdot w_j \cdot \sigma_j}{\sum_{k=1}^M w_k \cdot \sigma_k}$$

- Here,  $\sigma_j = \sqrt{\text{Var}(f(X) \mid X \in S_j)}$  is the standard deviation of the function within stratum  $S_j$ . This approach is optimal when the variance differs significantly across strata.

## Example: Estimating the Integral of a Function

Suppose we want to estimate the integral of a function  $f(x)$  over the interval  $[0, 1]$ :

$$I = \int_0^1 f(x) dx$$

Using stratified sampling, we can proceed as follows:

### 1. Divide the Interval:

- Split the interval  $[0, 1]$  into  $M$  equal subintervals (strata), say  $M = 4$ . Each stratum  $S_j$  is then an interval  $\left[\frac{j-1}{M}, \frac{j}{M}\right)$  for  $j = 1, 2, 3, 4$ .

### 2. Sample Within Each Stratum:

- Draw  $n_j$  samples uniformly within each subinterval. For simplicity, assume  $n_j = n/M$  samples are drawn from each stratum.

### 3. Estimate the Integral:

- Calculate the estimate within each stratum:

$$\hat{I}_j = \frac{1}{n_j} \sum_{i=1}^{n_j} f(X_{i,j})$$

- The overall estimate is the average of the stratum estimates, weighted by the size of each stratum (in this case, equal weights):

$$\hat{I}_{\text{Stratified}} = \frac{1}{M} \sum_{j=1}^M \hat{I}_j$$

By ensuring that samples are spread evenly across the interval, Stratified Sampling reduces the variance of the integral estimate, especially if  $f(x)$  varies significantly across the interval.

## Comparison to Standard Monte Carlo

- Standard Monte Carlo:** Samples are drawn uniformly from the entire interval, which could lead to clusters of samples in some regions and gaps in others. The estimator would be:

$$\hat{I}_{\text{MC}} = \frac{1}{n} \sum_{i=1}^n f(X_i)$$

with samples  $X_i$  drawn uniformly from  $[0, 1]$ .

- Stratified Sampling:** By dividing the interval into strata and ensuring each stratum is sampled, the variance of the estimate  $\hat{I}_{\text{Stratified}}$  is generally lower than that of  $\hat{I}_{\text{MC}}$ , leading to a more accurate estimate with the same number of samples.

## Benefits and Limitations

### • Benefits:

- Variance Reduction:** By ensuring even coverage of the sample space, Stratified Sampling typically results in lower variance compared to standard Monte Carlo, especially in cases where the function or distribution has significant variation across the space.
- Flexibility:** Stratified Sampling can be applied to a wide range of problems, from simple integrals to complex multi-dimensional simulations.

### • Limitations:

- Complexity:** Implementing Stratified Sampling requires partitioning the sample space and possibly adjusting the number of samples in each stratum, which adds complexity compared to standard Monte Carlo.

- **Choice of Strata:** The effectiveness of Stratified Sampling depends on how the strata are chosen. Poorly chosen strata may not lead to significant variance reduction.

## Latin Hypercube Sampling

### Latin Hypercube Sampling (LHS)

Latin Hypercube Sampling (LHS) is a sophisticated technique used in Monte Carlo simulations to generate samples in a way that ensures a more uniform coverage of the sample space. Unlike standard random sampling, which can lead to clustering or gaps in the sampled space, LHS systematically partitions the space into strata and samples from each stratum, which can significantly improve the efficiency and accuracy of simulations, especially in high-dimensional problems.

In particular, LHS reduces clustering and gaps in the sampled space, leading to more accurate estimates and reduced variance. While LHS is more complex to implement compared to standard random sampling, its benefits in terms of improved coverage and accuracy make it a valuable tool for a wide range of Monte Carlo applications, particularly in high-dimensional and complex simulations.

### Detailed Explanation of Latin Hypercube Sampling (LHS)

#### Basic Concept

In Latin Hypercube Sampling, the sample space is divided into equally probable intervals (strata) in each dimension. By sampling systematically within each stratum and ensuring that each stratum is represented, LHS reduces the chance of clustering and ensures that the sample points are spread more uniformly across the entire space. This is particularly useful for improving the accuracy of estimates and reducing the number of samples needed to achieve a given level of precision.

#### How it Works

##### 1. Partition the Sample Space:

- For each dimension  $i$  of the sample space, divide the interval  $[0, 1]$  into  $n$  equal intervals or strata. This results in  $n$  strata for each dimension.
- For example, if  $n = 5$ , the intervals for each dimension would be  $[0, 0.2)$ ,  $[0.2, 0.4)$ ,  $[0.4, 0.6)$ ,  $[0.6, 0.8)$ , and  $[0.8, 1)$ .

##### 2. Sample Within Each Stratum:

- Within each stratum, draw a sample from a uniform distribution. This means that for each stratum, you have one sample, but you need to ensure that each stratum is represented exactly once in each dimension.

##### 3. Create a Latin Hypercube:

- Construct a Latin Hypercube by arranging the samples such that each combination of strata from different dimensions appears exactly once. This means that if you have  $n$  dimensions, each sample will be drawn from a different stratum in each dimension.

#### 4. Generate the Final Sample Set:

- Randomly permute the samples within each stratum to ensure that the final sample set is uniformly distributed across the entire sample space.

### Detailed Steps for Generating LHS Samples

Let's consider a concrete example where we want to generate samples for a 2-dimensional space using LHS with  $n$  samples:

#### 1. Divide Each Dimension:

- Divide each dimension into  $n$  equal intervals. For instance, if  $n = 5$ , each dimension is divided into intervals  $[0, 0.2)$ ,  $[0.2, 0.4)$ ,  $[0.4, 0.6)$ ,  $[0.6, 0.8)$ , and  $[0.8, 1)$ .

#### 2. Draw Samples:

- Within each interval, draw one sample. For the first dimension, you might get samples  $[0.05, 0.25, 0.45, 0.65, 0.85]$  and for the second dimension, you might get samples  $[0.1, 0.3, 0.5, 0.7, 0.9]$ .

#### 3. Construct LHS Matrix:

- Create a matrix where each row corresponds to a sample, and each column corresponds to a dimension. Ensure that each row has a unique combination of strata from each dimension.

- For example, if you draw 5 samples, your matrix might look like this:

0.05	0.1
0.25	0.3
0.45	0.5
0.65	0.7
0.85	0.9

- Each combination of intervals (strata) is represented exactly once in the final matrix.

#### 4. Random Permutation:

- To further randomize the sample distribution, you can randomly permute the order of samples within each stratum. This ensures that the samples are uniformly distributed while still maintaining the Latin Hypercube property.

### Variance Reduction

The key advantage of LHS is its ability to cover the sample space more uniformly compared to simple random sampling. This uniform coverage helps to reduce the variance of the estimate because it ensures that each part of the sample space is represented, avoiding clustering and gaps.

The variance reduction can be understood by comparing the variance of the estimator obtained using LHS to that obtained using standard random sampling. LHS ensures that each interval is sampled exactly once, leading to more balanced representation across the entire sample space.



## Example: Estimating the Integral of a Function

Suppose we want to estimate the integral of a function  $f(x, y)$  over the unit square  $[0, 1] \times [0, 1]$ . Using LHS with  $n$  samples, the steps are:

### 1. Divide the Unit Square:

- Divide each dimension into  $n$  intervals. For instance, if  $n = 4$ , divide  $[0, 1]$  into intervals  $[0, 0.25)$ ,  $[0.25, 0.5)$ ,  $[0.5, 0.75)$ , and  $[0.75, 1)$ .

### 2. Draw Samples:

- Within each interval, draw one sample for each dimension.

### 3. Construct LHS Matrix:

- Arrange the samples such that each combination of intervals is represented exactly once.
- For example, you might end up with the following matrix:

0.1	0.2
0.3	0.5
0.6	0.8
0.8	0.9

### 4. Estimate the Integral:

- Compute the estimate of the integral using:

$$\hat{I}_{\text{LHS}} = \frac{1}{n} \sum_{i=1}^n f(X_i, Y_i)$$

By using LHS, the estimate of the integral benefits from reduced variance due to the more uniform coverage of the sample space.

## Benefits and Limitations

### • Benefits:

- **Uniform Coverage:** LHS provides a more uniform coverage of the sample space, reducing the chance of clustering and gaps, which improves the accuracy of estimates.
- **Variance Reduction:** By ensuring that each stratum is represented exactly once, LHS can reduce the variance of the Monte Carlo estimator, particularly in high-dimensional spaces.
- **Efficient Sampling:** For a given number of samples, LHS often requires fewer samples to achieve the same level of precision compared to standard random sampling.

### • Limitations:

- **Implementation Complexity:** LHS can be more complex to implement compared to standard random sampling, particularly for high-dimensional problems.
- **Stratum Selection:** The effectiveness of LHS depends on the choice of intervals (strata) and their division. Poor stratum selection may not lead to significant improvements.

## Applications

- **Engineering and Design:** In sensitivity analysis and optimization, LHS is used to sample design parameters more uniformly, leading to better insights into the behavior of complex systems.
- **Environmental Modeling:** LHS is used to simulate environmental processes, ensuring that different regions of the input space are adequately sampled.
- **Finance and Risk Analysis:** In financial simulations, LHS is used to model various scenarios and estimate risks more accurately.

## Quasi Monte Carlo

### Quasi Monte Carlo

Quasi-Monte Carlo (QMC) methods are advanced techniques used to generate sample points in a more structured manner than traditional Monte Carlo methods. Unlike random sampling, which relies on purely random points, QMC methods use deterministic sequences designed to cover the sample space more uniformly. This uniformity helps improve the accuracy and efficiency of simulations, particularly in high-dimensional problems.

By using low-discrepancy sequences, QMC methods reduce variance, improve accuracy, and achieve faster convergence rates in numerical simulations. Although QMC methods involve more complex implementation and may face limitations in very high dimensions, they offer significant advantages in terms of accuracy and efficiency for a wide range of applications, including numerical integration, financial modeling, engineering simulations, and scientific computing.

### Detailed Explanation of Quasi-Monte Carlo (QMC)

#### Basic Concept

Quasi-Monte Carlo methods aim to generate sequences of sample points that are distributed more uniformly over the sample space compared to random samples. The key idea is to use low-discrepancy sequences (also known as quasi-random sequences) instead of purely random sequences. These low-discrepancy sequences are designed to fill the sample space more uniformly, reducing the variance of the estimates and improving the convergence rate of the numerical integration.

#### Low-Discrepancy Sequences

Low-discrepancy sequences are deterministic sequences that provide more uniform coverage of the sample space. These sequences are also known as quasi-random sequences or quasi-random numbers. Common examples include:

##### 1. Van der Corput Sequence:

- A one-dimensional low-discrepancy sequence based on the binary expansion of the index. The sequence is generated by reversing the binary digits of the index and dividing by a power of 2.
- For example, the first few terms of the Van der Corput sequence in base 2 are: 0, 0.5, 0.25, 0.75, 0.125, 0.375, etc.

##### 2. Halton Sequence:

- A multi-dimensional extension of the Van der Corput sequence. Each dimension uses a different base to generate a sequence. For example, the first two dimensions might use bases 2 and 3.
- The Halton sequence for two dimensions can be constructed by combining the sequences generated for each base.

### 3. Sobol Sequence:

- A multi-dimensional low-discrepancy sequence that is widely used due to its excellent uniformity properties. The Sobol sequence is constructed using a set of generating vectors and operates in higher dimensions with high uniformity.

### 4. Niederreiter Sequence:

- Another multi-dimensional sequence that uses a different approach for generating low-discrepancy points. It is also known for its good coverage properties in high dimensions.

## How QMC Works

### 1. Generate Quasi-Random Sequence:

- Choose a low-discrepancy sequence appropriate for the problem's dimensionality. Generate  $n$  quasi-random points from this sequence. These points are deterministic and cover the sample space more uniformly than random points.

### 2. Apply to Numerical Integration:

- Use the quasi-random points to estimate integrals or other quantities. For example, if you want to estimate the integral of a function  $f(x)$  over a domain, you would use the quasi-random points to compute:

$$\hat{I}_{\text{QMC}} = \frac{1}{n} \sum_{i=1}^n f(x_i)$$

- Here,  $x_i$  are the quasi-random points generated by the sequence.

### 3. Evaluate Accuracy:

- Evaluate the accuracy of the estimate by comparing it with known values or using error bounds specific to the quasi-random sequence.

## Variance Reduction

QMC methods improve the efficiency of Monte Carlo simulations by providing more uniform coverage of the sample space. This uniformity reduces the variance of the estimator and improves convergence rates. Specifically:

- **Uniform Coverage:** Quasi-random sequences are designed to cover the sample space more uniformly than random samples, leading to more accurate estimates with fewer points.
- **Convergence Rate:** QMC methods typically exhibit convergence rates that are faster than random sampling. For example, the convergence rate for QMC methods is often proportional to  $\frac{1}{\sqrt{n}}$  for random sampling, but it can be significantly better for quasi-random sequences.

## Example: Estimating the Integral of a Function

Suppose we want to estimate the integral of a function  $f(x)$  over the unit interval  $[0, 1]$  using QMC:

### 1. Generate Quasi-Random Points:

- Use a low-discrepancy sequence, such as the Van der Corput sequence, to generate  $n$  quasi-random points in the interval  $[0, 1]$ .

### 2. Compute the Integral Estimate:

- Calculate the integral estimate using:

$$\hat{I}_{\text{QMC}} = \frac{1}{n} \sum_{i=1}^n f(x_i)$$

- Here,  $x_i$  are the quasi-random points.

### 3. Compare with Exact Value:

- Compare the estimate with the exact value of the integral to evaluate accuracy.

## Benefits and Limitations

### • Benefits:

- **Improved Accuracy:** QMC methods often provide more accurate estimates with fewer samples due to their uniform coverage of the sample space.
- **Faster Convergence:** The convergence rate of QMC methods is typically faster than that of standard Monte Carlo methods, especially in high-dimensional problems.
- **Reduced Variance:** By ensuring that the sample points are more uniformly distributed, QMC methods reduce the variance of the estimator.

### • Limitations:

- **Implementation Complexity:** Generating quasi-random sequences and implementing QMC methods can be more complex than standard Monte Carlo methods.
- **Dimensionality:** While QMC methods perform well in moderate dimensions, their effectiveness can decrease in very high dimensions, where the uniformity of coverage becomes less effective.
- **Problem-Specific Adaptation:** The choice of quasi-random sequence and its parameters may need to be adapted to the specific problem being solved.

## Multi dimensional Monte Carlo

### Multi-dimensional scenarios

Generating correlated geometric Brownian motion (GBM) paths involves several steps. For  $n$  GBM processes, each process  $i$  has its own parameters:

- $\mu_i$ : the drift or expected return
- $\sigma_i$ : the volatility or standard deviation
- $S_i(0)$ : the initial value of the process

- $T$ : the time horizon
- $N$ : the number of time steps

We can then define the pairwise correlations between each of the Brownian motions.

- Create an  $n \times n$  correlation matrix  $\mathbf{C}$ . This matrix should be positive definite.
- Perform the [Cholesky decomposition](Cholesky decomposition) on the correlation matrix  $\mathbf{C}$  to get a lower triangular matrix  $\mathbf{L}$  such that  $\mathbf{C} = \mathbf{L}\mathbf{L}^\top$ .  
This means that for each time step, simulate  $n$  independent standard normal variables  $Z_i(t)$ , for  $i = 1, 2, \dots, n$  and each time step  $t$ .

We then multiply the matrix  $\mathbf{L}$  by the vector of independent normal variables to obtain correlated normal variables:

$$W_i(t) = \sum_{j=1}^n L_{ij} Z_j(t)$$

- Here,  $W_i(t)$  represents the correlated Brownian motion increments.

The GBM for each process  $i$  can be constructed using the correlated Brownian motion paths:

$$S_i(t) = S_i(0) \exp \left( \left( \mu_i - \frac{\sigma_i^2}{2} \right) t + \sigma_i W_i(t) \right)$$

Where:

- $W_i(t)$  is the cumulative sum of the increments  $W_i(\Delta t)$  over time.
- $\Delta t = \frac{T}{N}$  is the time step size.  
Repeat the steps to simulate the GBM paths at all time steps from  $t = 0$  to  $t = T$ .

## Example: Python Pseudo-Code

Here's how the process might look in Python:

```

import numpy as np

# Parameters
n = 2 # Number of GBM processes
T = 1.0 # Time horizon
N = 100 # Number of time steps
mu = np.array([0.1, 0.2]) # Drift for each process
sigma = np.array([0.2, 0.3]) # Volatility for each process
S0 = np.array([100, 200]) # Initial values

# Time increment
dt = T / N

# Correlation matrix and Cholesky decomposition
C = np.array([[1.0, 0.8], [0.8, 1.0]])
L = np.linalg.cholesky(C)

# Simulate the paths
S = np.zeros((n, N+1))
S[:, 0] = S0

for t in range(1, N+1):
    Z = np.random.normal(0, 1, n)
    W = np.dot(L, Z) * np.sqrt(dt)
    S[:, t] = S[:, t-1] * np.exp((mu - 0.5 * sigma**2) * dt + sigma * W)

# S now contains the simulated GBM paths

```

This approach will give you  $n$  GBM paths that are correlated according to the matrix  $C$ . Each row in  $S$  represents the path of a different GBM process.

## Example: MATLAB Pseudo-Code

MATLAB

```
% Parameters
n = 2;                % Number of GBM processes
T = 1.0;              % Time horizon
N = 100;              % Number of time steps
mu = [0.1, 0.2];      % Drift for each process
sigma = [0.2, 0.3];    % Volatility for each process
S0 = [100, 200];      % Initial values

% Time increment
dt = T / N;

% Correlation matrix and Cholesky decomposition
C = [1.0, 0.8; 0.8, 1.0];
L = chol(C, 'lower'); % Cholesky decomposition to get lower triangular matrix

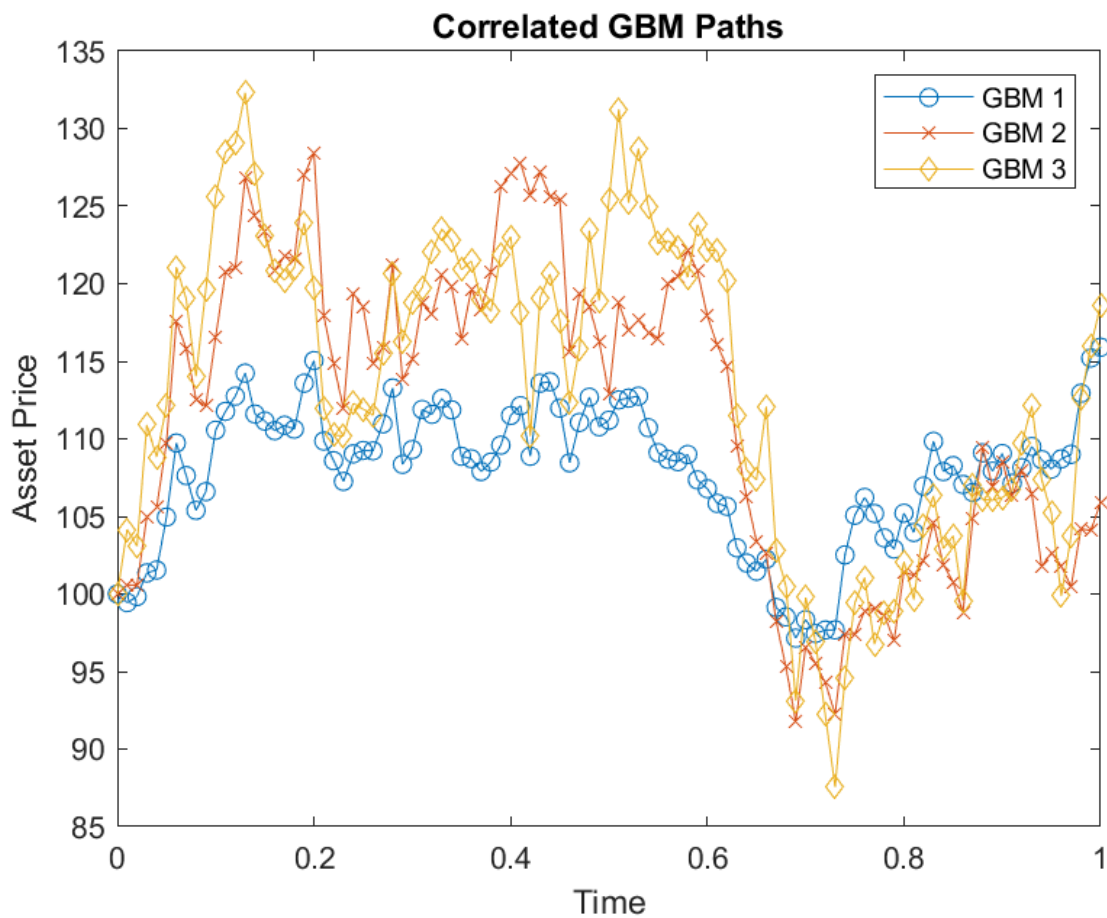
% Initialize matrix for storing GBM paths
S = zeros(n, N+1);
S(:, 1) = S0; % Set initial values

% Simulate the GBM paths
for t = 2:N+1
    Z = randn(n, 1); % Generate independent standard normal random variables
    W = L * Z * sqrt(dt); % Generate correlated Brownian increments
    S(:, t) = S(:, t-1) .* exp((mu' - 0.5 * sigma' .* sigma) * dt + sigma' .* W);
end

% Plot the results
time = linspace(0, T, N+1);
figure;
plot(time, S(1,:), '-o', time, S(2,:), '-x');
xlabel('Time');
ylabel('Asset Price');
legend('GBM 1', 'GBM 2');
title('Correlated GBM Paths');
```

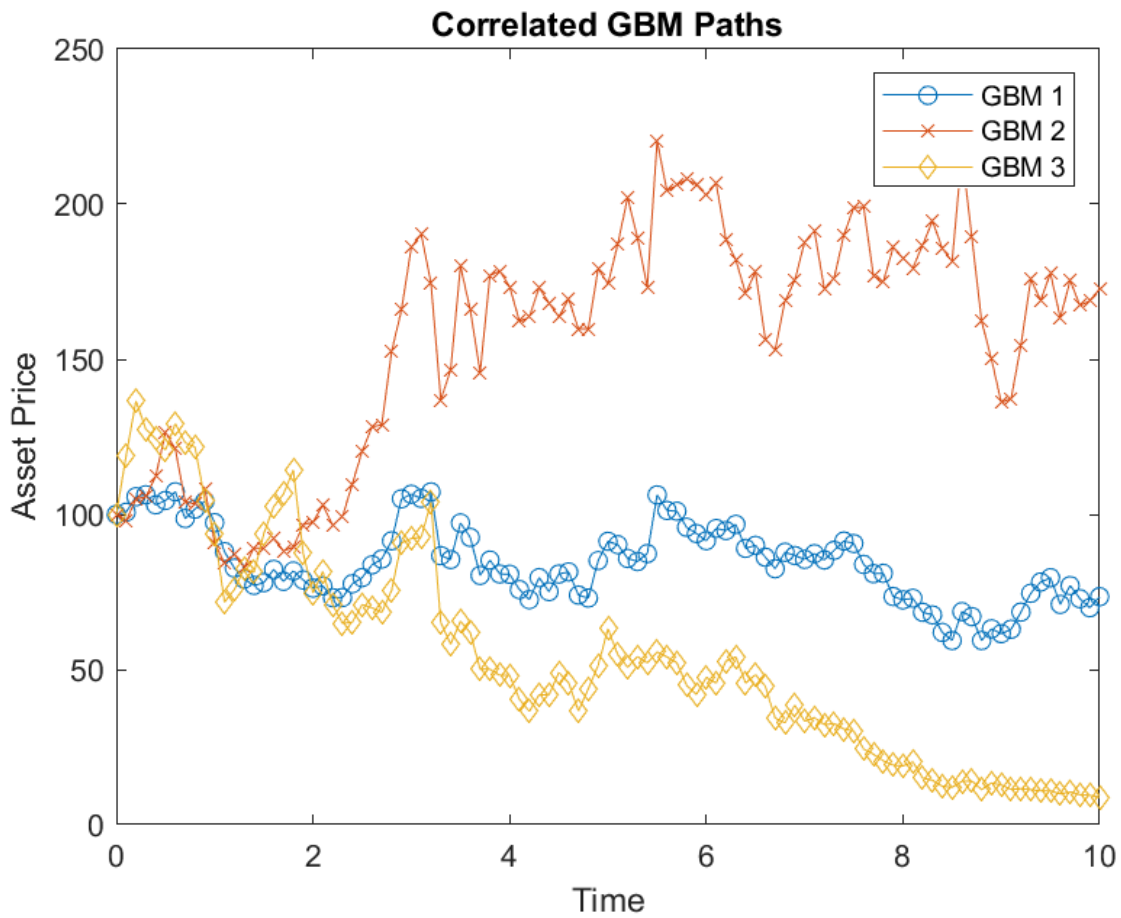
Below is an example of 3 correlated random paths. The correlation matrix is given by

$$\mathbf{C} = \begin{bmatrix} 1.0 & 0.8 & 0.7 \\ 0.8 & 1.0 & 0.5 \\ 0.7 & 0.5 & 1.0 \end{bmatrix}$$




We can clearly observe the effect of the positive correlation on the same paths. The 3 price paths start together and stay more or less together. However, due to the nature of these price paths, this is not necessarily true over longer horizons. When we change the time horizon to 10 years and generate 3 sample paths, they can easily diverge but remain correlated as can be seen in the graph below:





The effect is most clearly seen between  $GBM_1$  and  $GBM_2$  that have an 80% correlation and indeed the timing of the peaks coincides (visually) but the trajectories clearly diverge.

One might think this can be controlled by using a very high correlation, but even if we use a correlation of 99.99% between each of the brownian motions, we still get this effect



Using [Cointegration](#) rather than correlation can help mitigate this problem although it is not always easy to estimate the cointegration between two time series.

## Cointegration

### Overview of Cointegration

Cointegration is a concept in time series analysis that refers to a long-run equilibrium relationship between two or more non-stationary time series. Although these series may individually wander over time due to stochastic trends, they move together in such a way that a certain linear combination of them is stationary, meaning it has a constant mean and variance over time.

If non-stationary series are not cointegrated, regressions between them can lead to misleading results, known as spurious regression. Cointegration relationships may change over time due to structural breaks in the data, requiring techniques like regime-switching models.

### Stationarity and Non-Stationarity

Before diving into cointegration, it is important to understand stationarity and non-stationarity:

- **Stationary Time Series:** A time series is said to be stationary if its statistical properties, such as mean, variance, and autocorrelation, do not change over time. Formally, a time series  $y_t$  is stationary if for all time points  $t$ , the joint probability distribution of  $y_t, y_{t+k}, \dots$  does not depend on  $t$ .
- **Non-Stationary Time Series:** A time series is non-stationary if its statistical properties change over time. A common type of non-stationarity is a unit root process, where the time

series exhibits a stochastic trend.

One of the most common tests for stationarity is the **Augmented Dickey-Fuller (ADF) test**, which tests the null hypothesis that a unit root is present in a time series sample.

## Unit Roots and Integration Order

A time series  $y_t$  is said to be integrated of order  $d$ , denoted  $y_t \sim I(d)$ , if it becomes stationary after differencing  $d$  times:

- $I(0)$ : The series is stationary.
- $I(1)$ : The series has to be differenced once to become stationary (e.g., a random walk).
- $I(2)$ : The series has to be differenced twice to become stationary, and so on.

## Cointegration Concept

Consider two non-stationary time series  $y_t$  and  $x_t$ , both integrated of order 1, i.e.,  $y_t \sim I(1)$  and  $x_t \sim I(1)$ . Individually, these series are non-stationary and could drift apart over time. However, if there exists a linear combination of these series such that:

$$z_t = y_t - \beta x_t$$

where  $\beta$  is a constant, and  $z_t$  is stationary ( $z_t \sim I(0)$ ), then  $y_t$  and  $x_t$  are said to be **cointegrated**. The parameter  $\beta$  is known as the **cointegrating coefficient**.

In general, for a set of  $n$  time series  $\mathbf{y}_t = (y_{1,t}, y_{2,t}, \dots, y_{n,t})'$ , if there exists a vector  $\beta = (\beta_1, \beta_2, \dots, \beta_n)'$  such that:

$$\mathbf{z}_t = \beta' \mathbf{y}_t = \beta_1 y_{1,t} + \beta_2 y_{2,t} + \dots + \beta_n y_{n,t} \sim I(0)$$

then the series are cointegrated, and  $\beta$  is called the **cointegrating vector**.

## Economic Interpretation

Cointegration has significant implications in economics and finance. For instance, if two financial assets are cointegrated, this suggests a long-term equilibrium relationship between their prices. Even though each asset may be non-stationary on its own, their prices are linked over time, possibly due to arbitrage opportunities or underlying economic relationships.

## Engle-Granger Two-Step Method

The most straightforward approach to testing for cointegration between two time series is the **Engle-Granger Two-Step Method**:

### 1. Step 1: Estimate the Cointegrating Relationship

Regress one series on the other using Ordinary Least Squares (OLS) to estimate the cointegrating coefficient:

$$y_t = \alpha + \beta x_t + \epsilon_t$$

where  $\epsilon_t$  is the residual term.

### 2. Step 2: Test the Residuals for Stationarity

Apply a unit root test (such as the ADF test) to the residuals  $\epsilon_t$ . If the residuals are stationary, i.e.,  $\epsilon_t \sim I(0)$ , then  $y_t$  and  $x_t$  are cointegrated.

## Johansen Cointegration Test

The Engle-Granger method is limited to testing for a single cointegrating relationship. The **Johansen Cointegration Test** extends this to multiple series, allowing for the identification of more than one cointegrating vector. The test is based on a Vector Autoregressive (VAR) framework.

### 1. Vector Autoregression (VAR)

Consider a VAR model of  $k$  non-stationary series  $\mathbf{y}_t$ :

$$\Delta \mathbf{y}_t = \mathbf{\Pi} \mathbf{y}_{t-1} + \sum_{i=1}^{k-1} \mathbf{\Gamma}_i \Delta \mathbf{y}_{t-i} + \mathbf{u}_t$$

where:

- $\Delta \mathbf{y}_t$  is the first difference of  $\mathbf{y}_t$ .
- $\mathbf{\Pi}$  is the long-run impact matrix.
- $\mathbf{\Gamma}_i$  are short-run impact matrices.
- $\mathbf{u}_t$  is a vector of error terms.

### 2. Eigenvalue Decomposition

The Johansen test involves decomposing the matrix  $\mathbf{\Pi}$  into:

$$\mathbf{\Pi} = \alpha \beta'$$

where:

- $\alpha$  represents the adjustment coefficients.
- $\beta$  represents the cointegrating vectors.

The rank of  $\mathbf{\Pi}$  determines the number of cointegrating vectors. If the rank is  $r$ , then there are  $r$  linearly independent cointegrating vectors.

### 3. Test Statistics

The Johansen test provides two test statistics:

#### 1. Trace Test:

$$\text{Trace Statistic} = -T \sum_{i=r+1}^n \ln(1 - \lambda_i)$$

where  $\lambda_i$  are the eigenvalues of  $\mathbf{\Pi}$ ,  $T$  is the sample size, and  $r$  is the number of cointegrating vectors.

#### 2. Maximum Eigenvalue Test:

$$\text{Max Eigenvalue Statistic} = -T \ln(1 - \lambda_{r+1})$$

## Error Correction Model (ECM)

When two or more time series are cointegrated, deviations from the long-run equilibrium will be corrected over time. This can be captured by an **Error Correction Model (ECM)**, which combines short-run dynamics with the long-run equilibrium:

$$\Delta y_t = \alpha(\beta x_t - y_{t-1}) + \gamma \Delta x_t + \epsilon_t$$

where:

- $\alpha$  is the speed of adjustment parameter.
- $\beta x_t - y_{t-1}$  represents the error correction term, which corrects deviations from the long-run equilibrium.

## Applications of Cointegration

Cointegration is widely used in various fields, including:

- **Economics:** To model long-term relationships between variables like GDP, interest rates, and inflation.
- **Finance:** To identify pairs of assets for statistical arbitrage (pairs trading).
- **Macroeconomics:** To analyze the relationship between consumption, income, and wealth.

## Example: Modeling Power and Gas Prices Using Cointegration

Let's walk through an example where we use cointegration to model the relationship between power and gas prices. Power and gas prices often exhibit a strong relationship because natural gas is a key input in electricity generation. Despite being individually non-stationary, their prices might move together in the long run, making them potential candidates for cointegration.

### Step 1: Understanding the Data

Suppose we have daily time series data for natural gas prices ( $P_{\text{gas}}$ ) and electricity prices ( $P_{\text{power}}$ ) over a period of several years. Both price series are likely non-stationary due to underlying trends, seasonality, and other factors.

Let's denote:

- $P_{\text{gas},t}$ : Natural gas price at time  $t$ .
- $P_{\text{power},t}$ : Electricity price at time  $t$ .

### Step 2: Test for Stationarity

We begin by testing each series for stationarity using the Augmented Dickey-Fuller (ADF) test.

- **For gas prices:**

ADF Test on  $P_{\text{gas},t}$  : Null hypothesis  $H_0$  : Series has a unit root (non-stationary).

- **For power prices:**

ADF Test on  $P_{\text{power},t}$  : Null hypothesis  $H_0$  : Series has a unit root (non-stationary).

If both tests fail to reject the null hypothesis, then both series are non-stationary, likely integrated of order 1,  $I(1)$ .

### Step 3: Test for Cointegration

Next, we check whether there is a long-run equilibrium relationship between the two series, i.e., whether they are cointegrated.

- We perform an Ordinary Least Squares (OLS) regression of power prices on gas prices:

$$P_{\text{power},t} = \alpha + \beta P_{\text{gas},t} + \epsilon_t$$

where  $\epsilon_t$  is the residual.

- **Residual Analysis:**

$$\epsilon_t = P_{\text{power},t} - \alpha - \beta P_{\text{gas},t}$$

We then test the residuals  $\epsilon_t$  for stationarity using the ADF test:

ADF Test on  $\epsilon_t$  : Null hypothesis  $H_0$  : Residuals have a unit root (non-stationary).

If the ADF test rejects the null hypothesis, the residuals are stationary, indicating that power and gas prices are cointegrated.

### Step 4: Estimating the Cointegrating Relationship

Suppose we find that power and gas prices are cointegrated. The OLS regression gives us the cointegrating relationship:

$$P_{\text{power},t} = \alpha + \beta P_{\text{gas},t} + \epsilon_t$$

The coefficient  $\beta$  represents the long-term equilibrium relationship between power and gas prices. For instance, if  $\beta = 2$ , it suggests that for every \$1 increase in the price of natural gas, the electricity price increases by \$2 in the long run.

### Step 5: Modeling with Error Correction

Given that the series are cointegrated, we can use an Error Correction Model (ECM) to capture both the short-term dynamics and the long-term relationship:

$$\Delta P_{\text{power},t} = \alpha_0 + \alpha_1(\epsilon_{t-1}) + \gamma_1 \Delta P_{\text{gas},t} + u_t$$

Where:

- $\Delta P_{\text{power},t}$  is the change in electricity prices.
- $\epsilon_{t-1}$  is the lagged residual from the cointegration equation, representing the error correction term.
- $\gamma_1 \Delta P_{\text{gas},t}$  captures the short-term impact of changes in gas prices on power prices.
- $\alpha_1$  is the speed of adjustment coefficient, showing how quickly deviations from the long-term equilibrium are corrected.

### Step 6: Interpretation

- **Long-Term Relationship:** The cointegrating equation  $P_{\text{power},t} = \alpha + \beta P_{\text{gas},t} + \epsilon_t$  suggests that power prices and gas prices move together in the long run.
- **Short-Term Dynamics:** The ECM indicates how short-term changes in gas prices affect power prices, and how power prices adjust to deviations from the long-term equilibrium.

## Step 7: Model Validation and Application

Finally, we would validate the model by checking the residuals of the ECM for autocorrelation and other issues. Once validated, this model can be used for forecasting power prices based on expected movements in gas prices, considering both short-term fluctuations and long-term equilibrium.

## Least Squares Monte Carlo (LSMC)

The Least Squares Monte Carlo (LSMC) method, introduced by Longstaff and Schwartz in 2001, is a powerful technique for valuing American options. This method combines Monte Carlo simulation with regression analysis to handle the early exercise feature of American options.

It provides a flexible approach to handle the early exercise feature of American options, and it can be applied to a wide range of financial derivatives with different complexities. The method involves simulating asset paths, estimating continuation values using regression, and back-calculating the option value through averaging discounted payoffs.

### Least Squares Monte Carlo (LSMC) Overview

The purpose of the LSMC method is used to estimate the value of American options by simulating multiple paths of the underlying asset and using regression analysis to determine the optimal exercise strategy. It is a direct extension of the general Monte Carlo method. This means it breaks down in similar steps and it also implies that all control variate techniques can be used.

As a first step, we will generate a large number of paths for the underlying asset price using Monte Carlo simulation. Second, at each time step before maturity, estimate the **continuation value** of the option using regression on the simulated paths.

By comparing the immediate exercise value with the estimated continuation value, we can decide whether to exercise the option early or continue holding it. Similarly to the Monte Carlo method, the average of the discounted payoffs along the simulated paths then determine the option value.

### Detailed Steps and Formulas

#### Simulate Asset Paths

Simulate paths of the underlying asset price using a stochastic process, typically geometric Brownian motion. The asset price  $S_t$  follows:

$$S_{t+\Delta t} = S_t \cdot \exp \left( (r - 0.5\sigma^2)\Delta t + \sigma\sqrt{\Delta t}Z_t \right)$$

where  $Z_t$  is a standard normal random variable.

#### Estimate Continuation Value

For each path, calculate the continuation value of the option at each time step using regression. This involves:

1. **Regression Analysis:**

- At each time step  $t$ , where the option is in the money (i.e.,  $S_t < K$ ), perform a regression to estimate the continuation value. The continuation value  $C_t$  is modeled as a function of the state variables (e.g., the asset price).
- Fit a regression model:

$$C_t = \beta_0 + \beta_1 S_t + \beta_2 S_t^2$$

where  $\beta_0, \beta_1, \beta_2$  are the regression coefficients.

## 2. Determine Optimal Exercise:

- Compare the payoff from immediate exercise with the continuation value:

$$\text{Payoff} = \max(K - S_t, 0)$$

- The option value at time  $t$  is the maximum of the immediate payoff and the continuation value:

$$\text{Value} = \max(\text{Payoff}, C_t)$$

## Back-Calculate the Option Value

Discount the payoffs from the final time step to the present value. For each path, calculate the average discounted payoff:

$$\text{Option Value} = \frac{1}{N} \sum_{i=1}^N \text{Discounted Payoff}_i$$

where  $N$  is the number of simulated paths.



## Example and Code

### Matlab Example

MATLAB

```
% Parameters
S0 = 100; K = 100; r = 0.05; sigma = 0.2; T = 1; N = 1000; M = 100;
dt = T / M; Smax = 200;
S = linspace(0, Smax, N+1);
payoffs = max(K - S, 0);

% Simulate asset paths
paths = zeros(N+1, M+1);
paths(:,1) = S0;
for i = 2:M+1
    Z = randn(N, 1);
    paths(:,i) = paths(:,i-1) .* exp((r - 0.5 * sigma^2) * dt + sigma *
sqrt(dt) * Z);
end

% Initialize option values at maturity
V = max(K - paths(:,end), 0);

% Backward induction
for i = M:-1:1
    in_the_money = paths(:,i) < K;
    X = paths(in_the_money, i);
    Y = V(in_the_money);

    % Perform regression
    X_poly = [ones(size(X)) X X.^2];
    beta = (X_poly' * X_poly) \ (X_poly' * Y);
    continuation_value = [ones(size(paths(:,i))) paths(:,i) paths(:,i).^2]
* beta;

    % Determine option value
    V = max(K - paths(:,i), exp(-r * dt) * continuation_value);
end

% Option value at time 0
option_price = mean(V);
disp(['American Put Option Price: ', num2str(option_price)]);
```

## Python Example

PYTHON

```
import numpy as np
from sklearn.linear_model import LinearRegression

# Parameters
S0 = 100; K = 100; r = 0.05; sigma = 0.2; T = 1; N = 1000; M = 100;
dt = T / M; Smax = 200
S = np.linspace(0, Smax, N+1)
payoffs = np.maximum(K - S, 0)

# Simulate asset paths
paths = np.zeros((N+1, M+1))
paths[:,0] = S0
for i in range(1, M+1):
    Z = np.random.normal(0, 1, N)
    paths[:,i] = paths[:,i-1] * np.exp((r - 0.5 * sigma**2) * dt + sigma *
np.sqrt(dt) * Z)

# Initialize option values at maturity
V = np.maximum(K - paths[:, -1], 0)

# Backward induction
for i in range(M-1, -1, -1):
    in_the_money = paths[:,i] < K
    X = paths[in_the_money, i].reshape(-1, 1)
    Y = V[in_the_money]

    # Perform regression
    X_poly = np.hstack((np.ones_like(X), X, X**2))
    reg = LinearRegression().fit(X_poly, Y)
    continuation_value =
reg.predict(np.hstack((np.ones_like(paths[:,i].reshape(-1,1)),
paths[:,i].reshape(-1,1), paths[:,i].reshape(-1,1)**2)))

    # Determine option value
    V = np.maximum(K - paths[:,i], np.exp(-r * dt) * continuation_value)

# Option value at time 0
option_price = np.mean(V)
print(f'American Put Option Price: {option_price}')
```